

Hierarchical KV Cache Compression for LLMs

2024324064 Joshua Dela Rosa

2025313067 Mengjie Li

2025-12-03 (수)

Contents

- Problem Statement
- Motivation
- Implementation
- Experimental Setup
- Preliminary Results and Discussions
- Conclusions

Transformers and Attention

- Transformers rely heavily in the attention mechanism
- Compares current token to all tokens in the sequence.

The pizza came out of the oven and **it** tasted good.

Strong connection between “pizza” and “it”

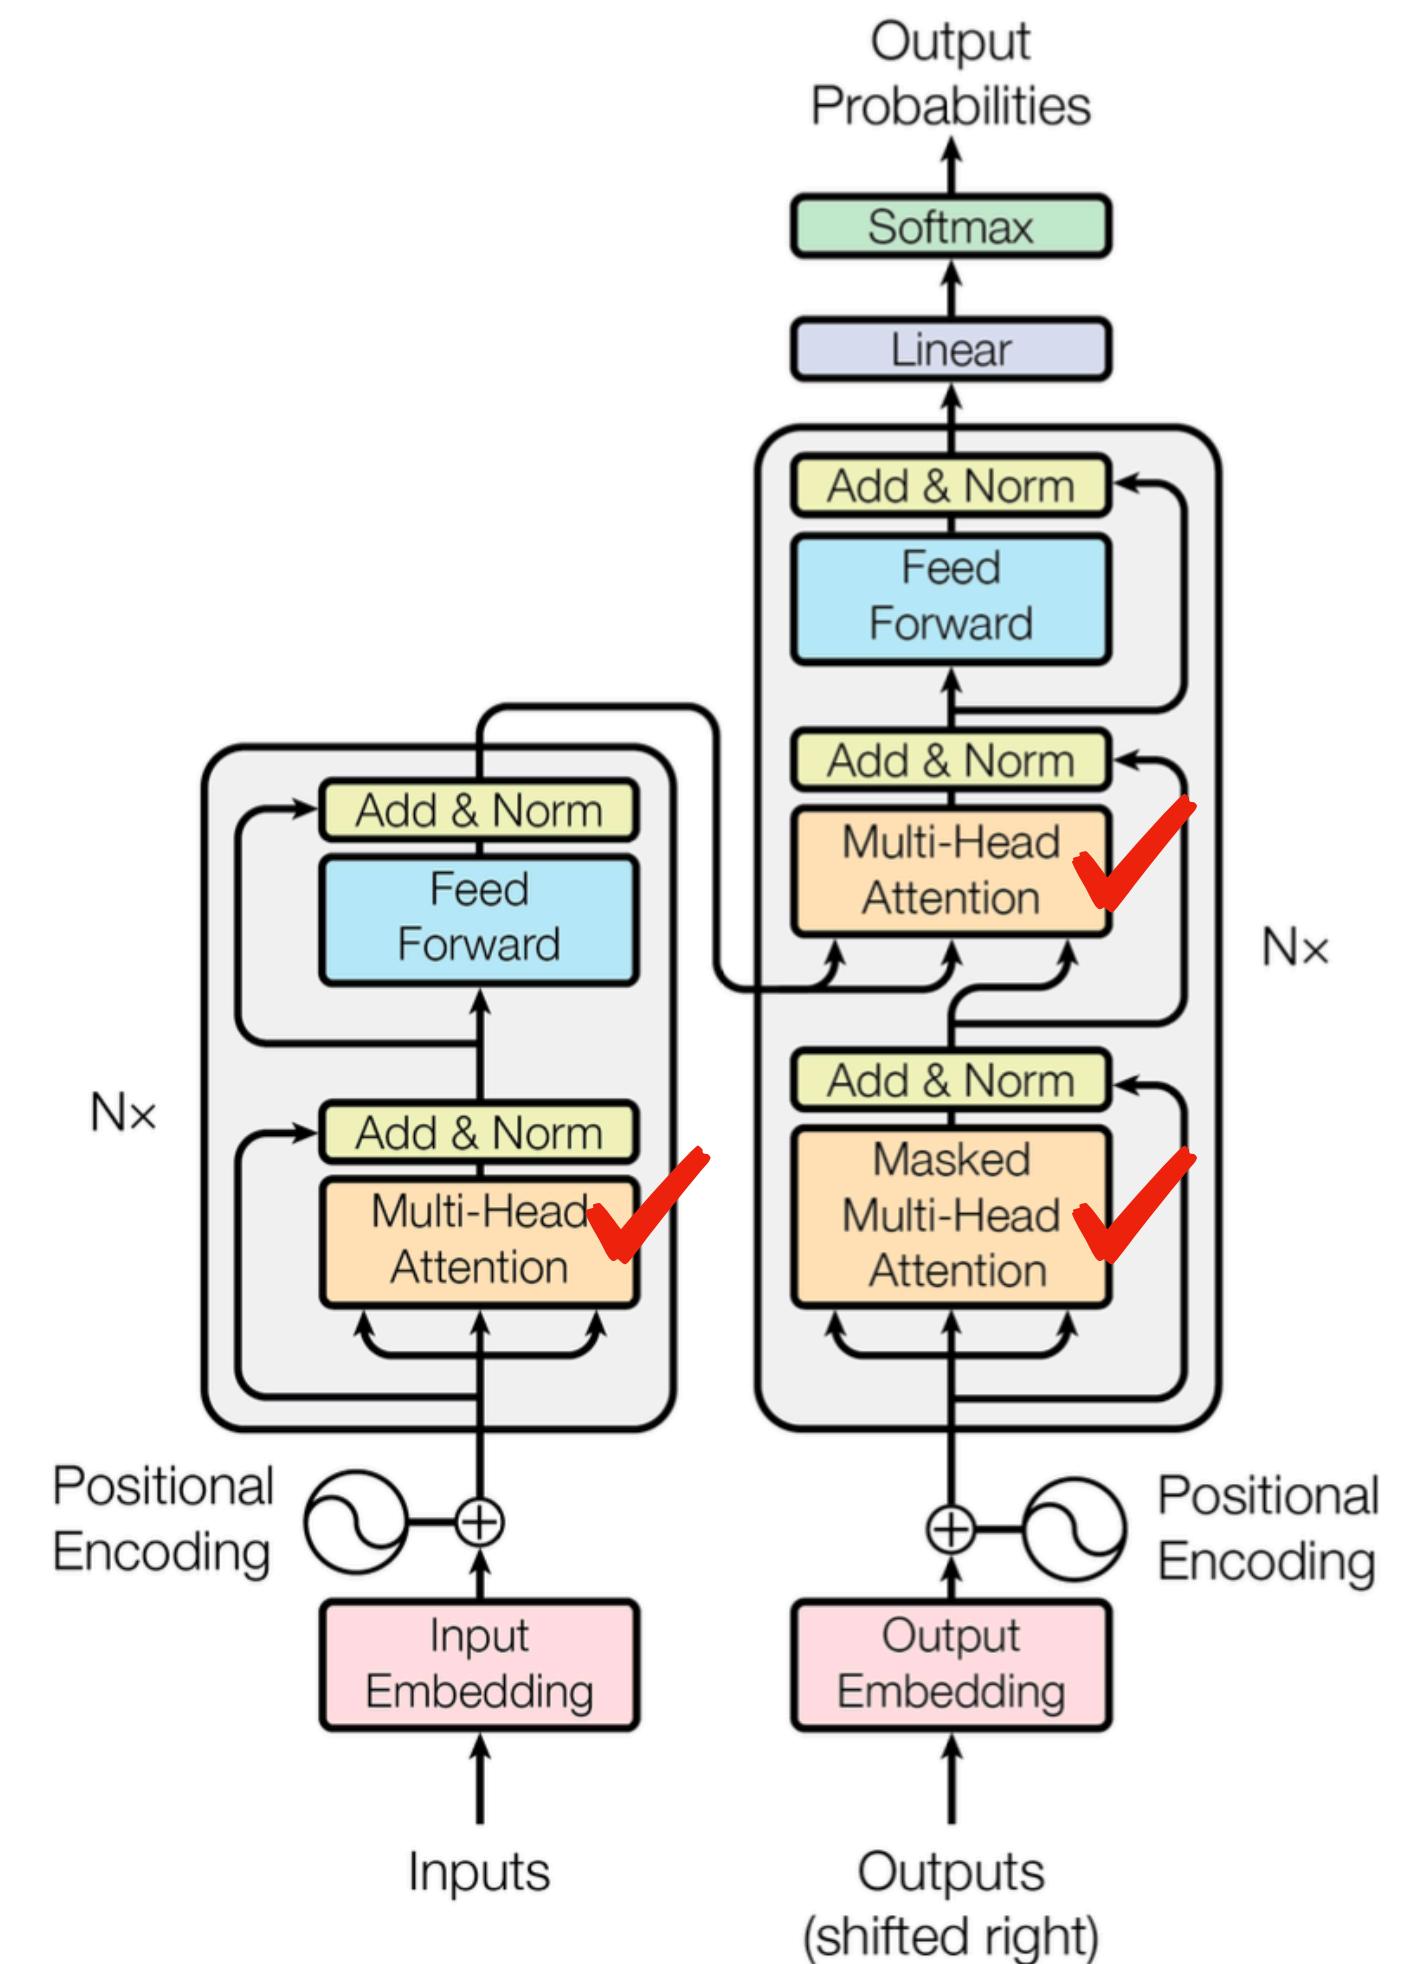
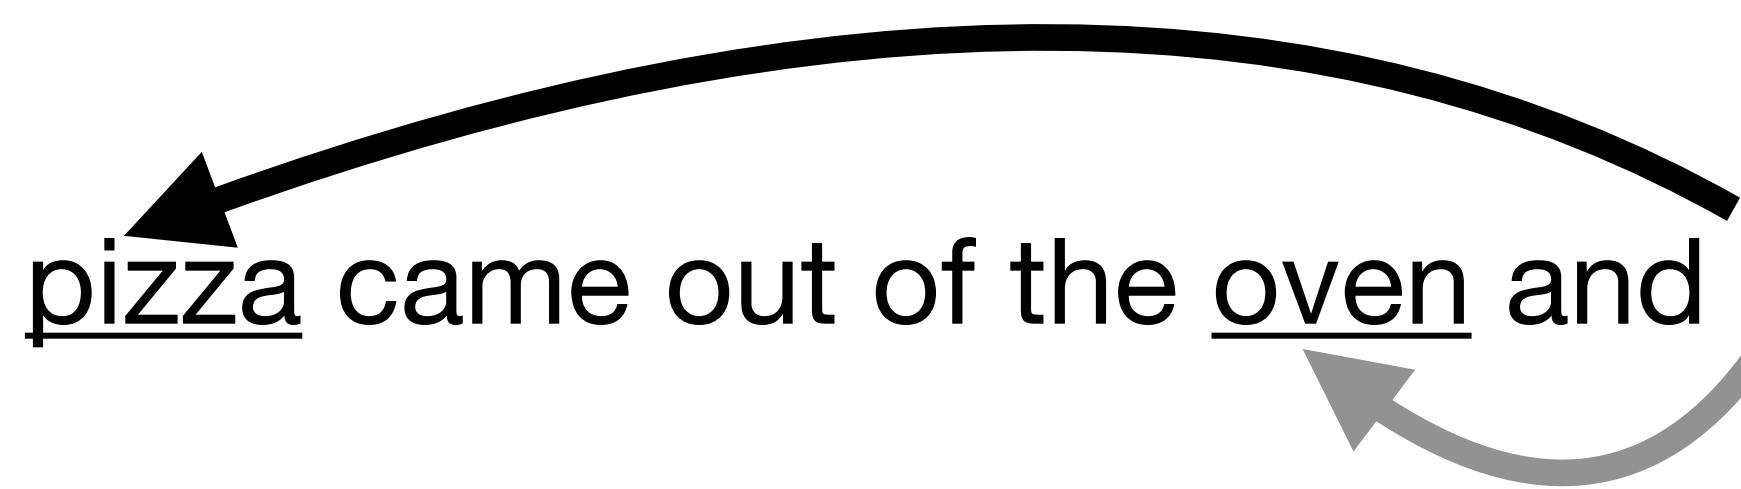
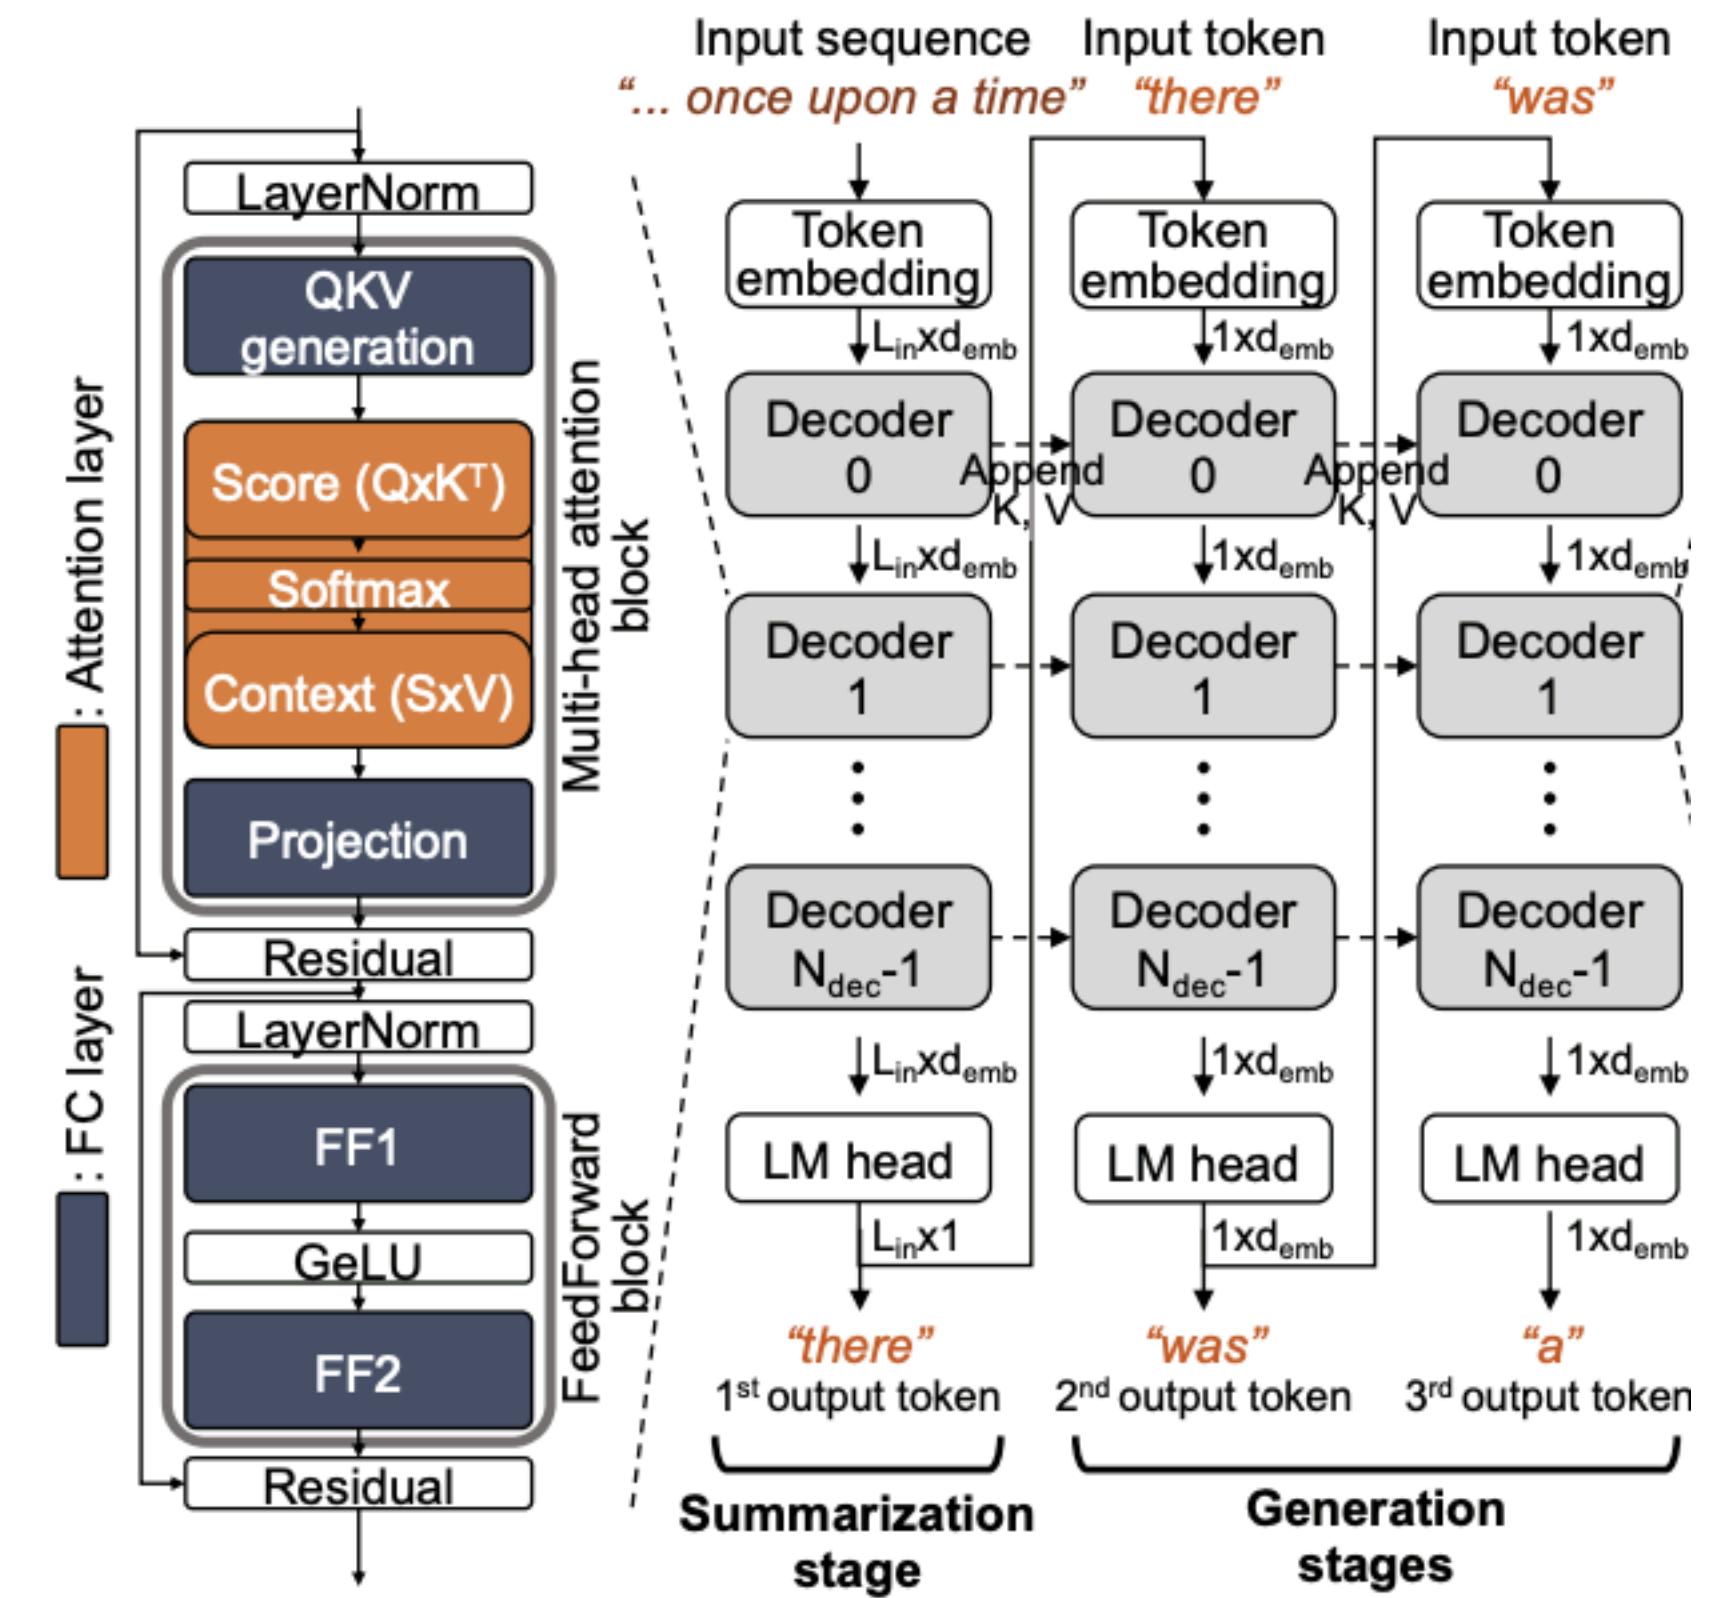


Figure 1: The Transformer - model architecture.

Transformers and Attention

- LLM computation is divided into two phases:
 - **Summarization (Prefill) Stage:** processing of all the tokens in the prompt
 - **Generation (Decode) Stage:** autoregressive generation of tokens

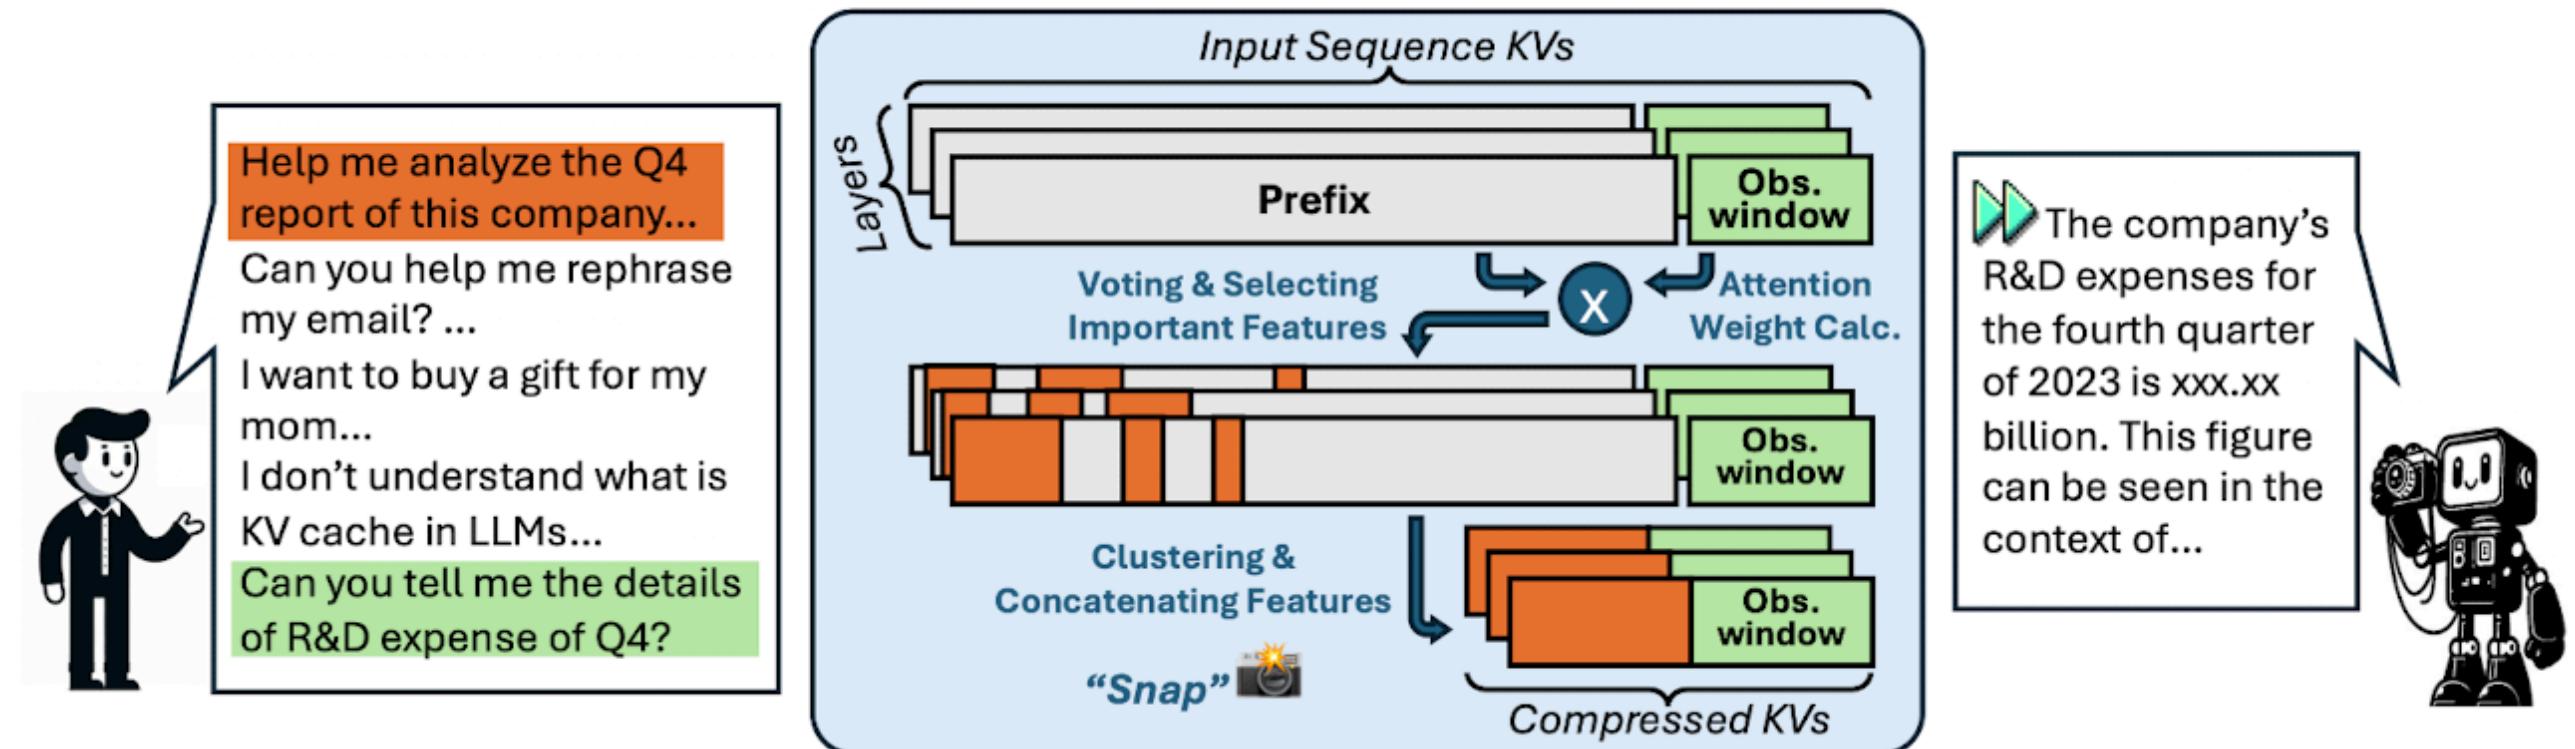


KV Cache

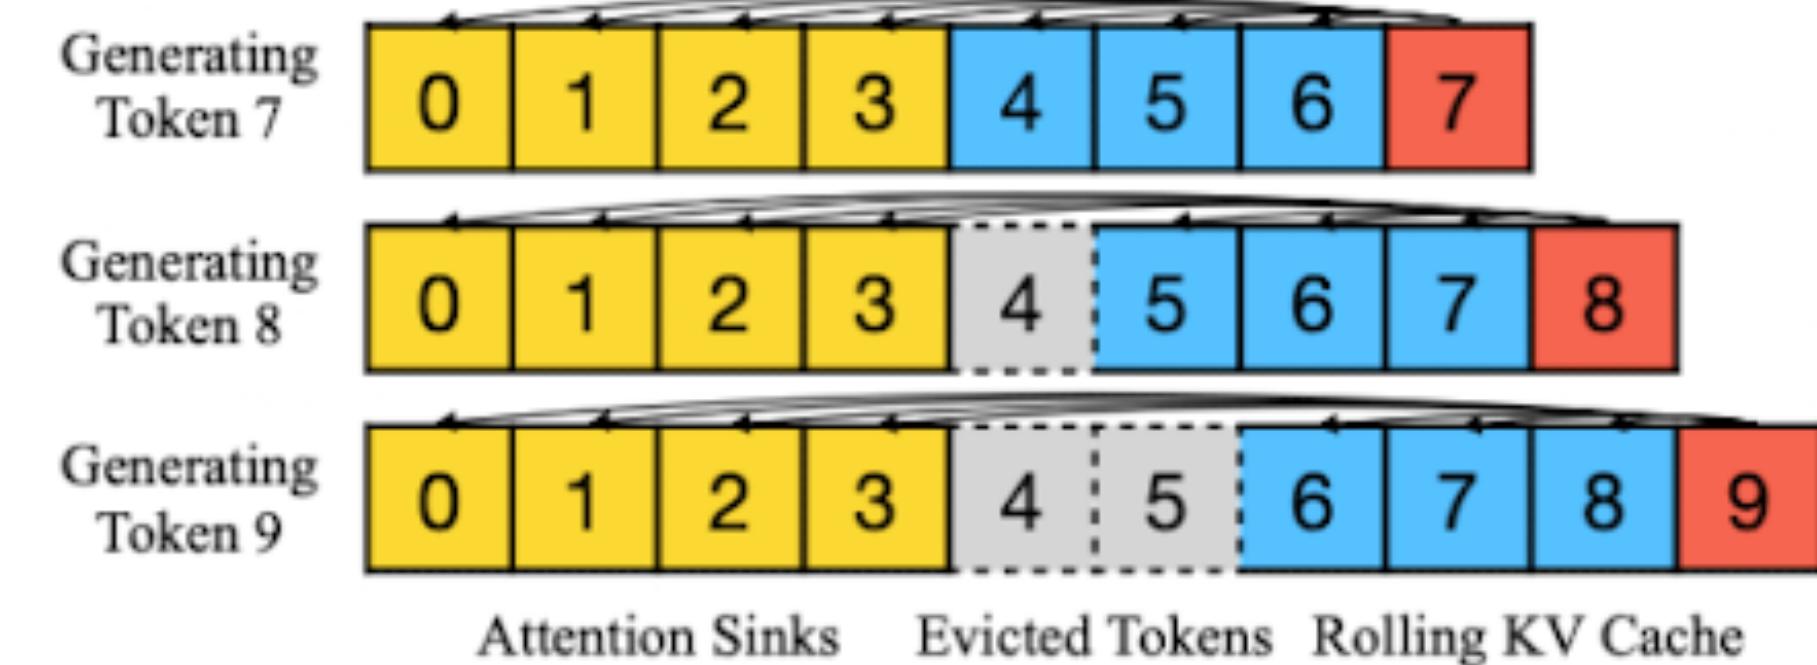
- To generate a new token, it needs the computed keys and values for all previous tokens.
- For computation efficiency, cache the KV, trade for extra memory.
- Problem:
 - longer context and larger LLMs = larger KV cache
- KV cache compression is a hot topic.
 - Token-level KV cache optimization

Current KV Cache Research

- **SnapKV (NIPS 2024)**
 - the end part of the prompt dictates which tokens of the prompt should be kept



- **AttentionSink (ICLR 2024)**
 - it was observed that a large amount of attention scores are allocated to the initial tokens



Problems with SnapKV

- Why always the end part of the prompt?
 - Suppose a prompt where the question is at the front.

Example Prompt A

Today was ...
For some reason **he fell sick** ...
So that was ...

Why was John absent yesterday?

- SnapKV is limited by the assumption that the most important parts are at the end of the prompt.

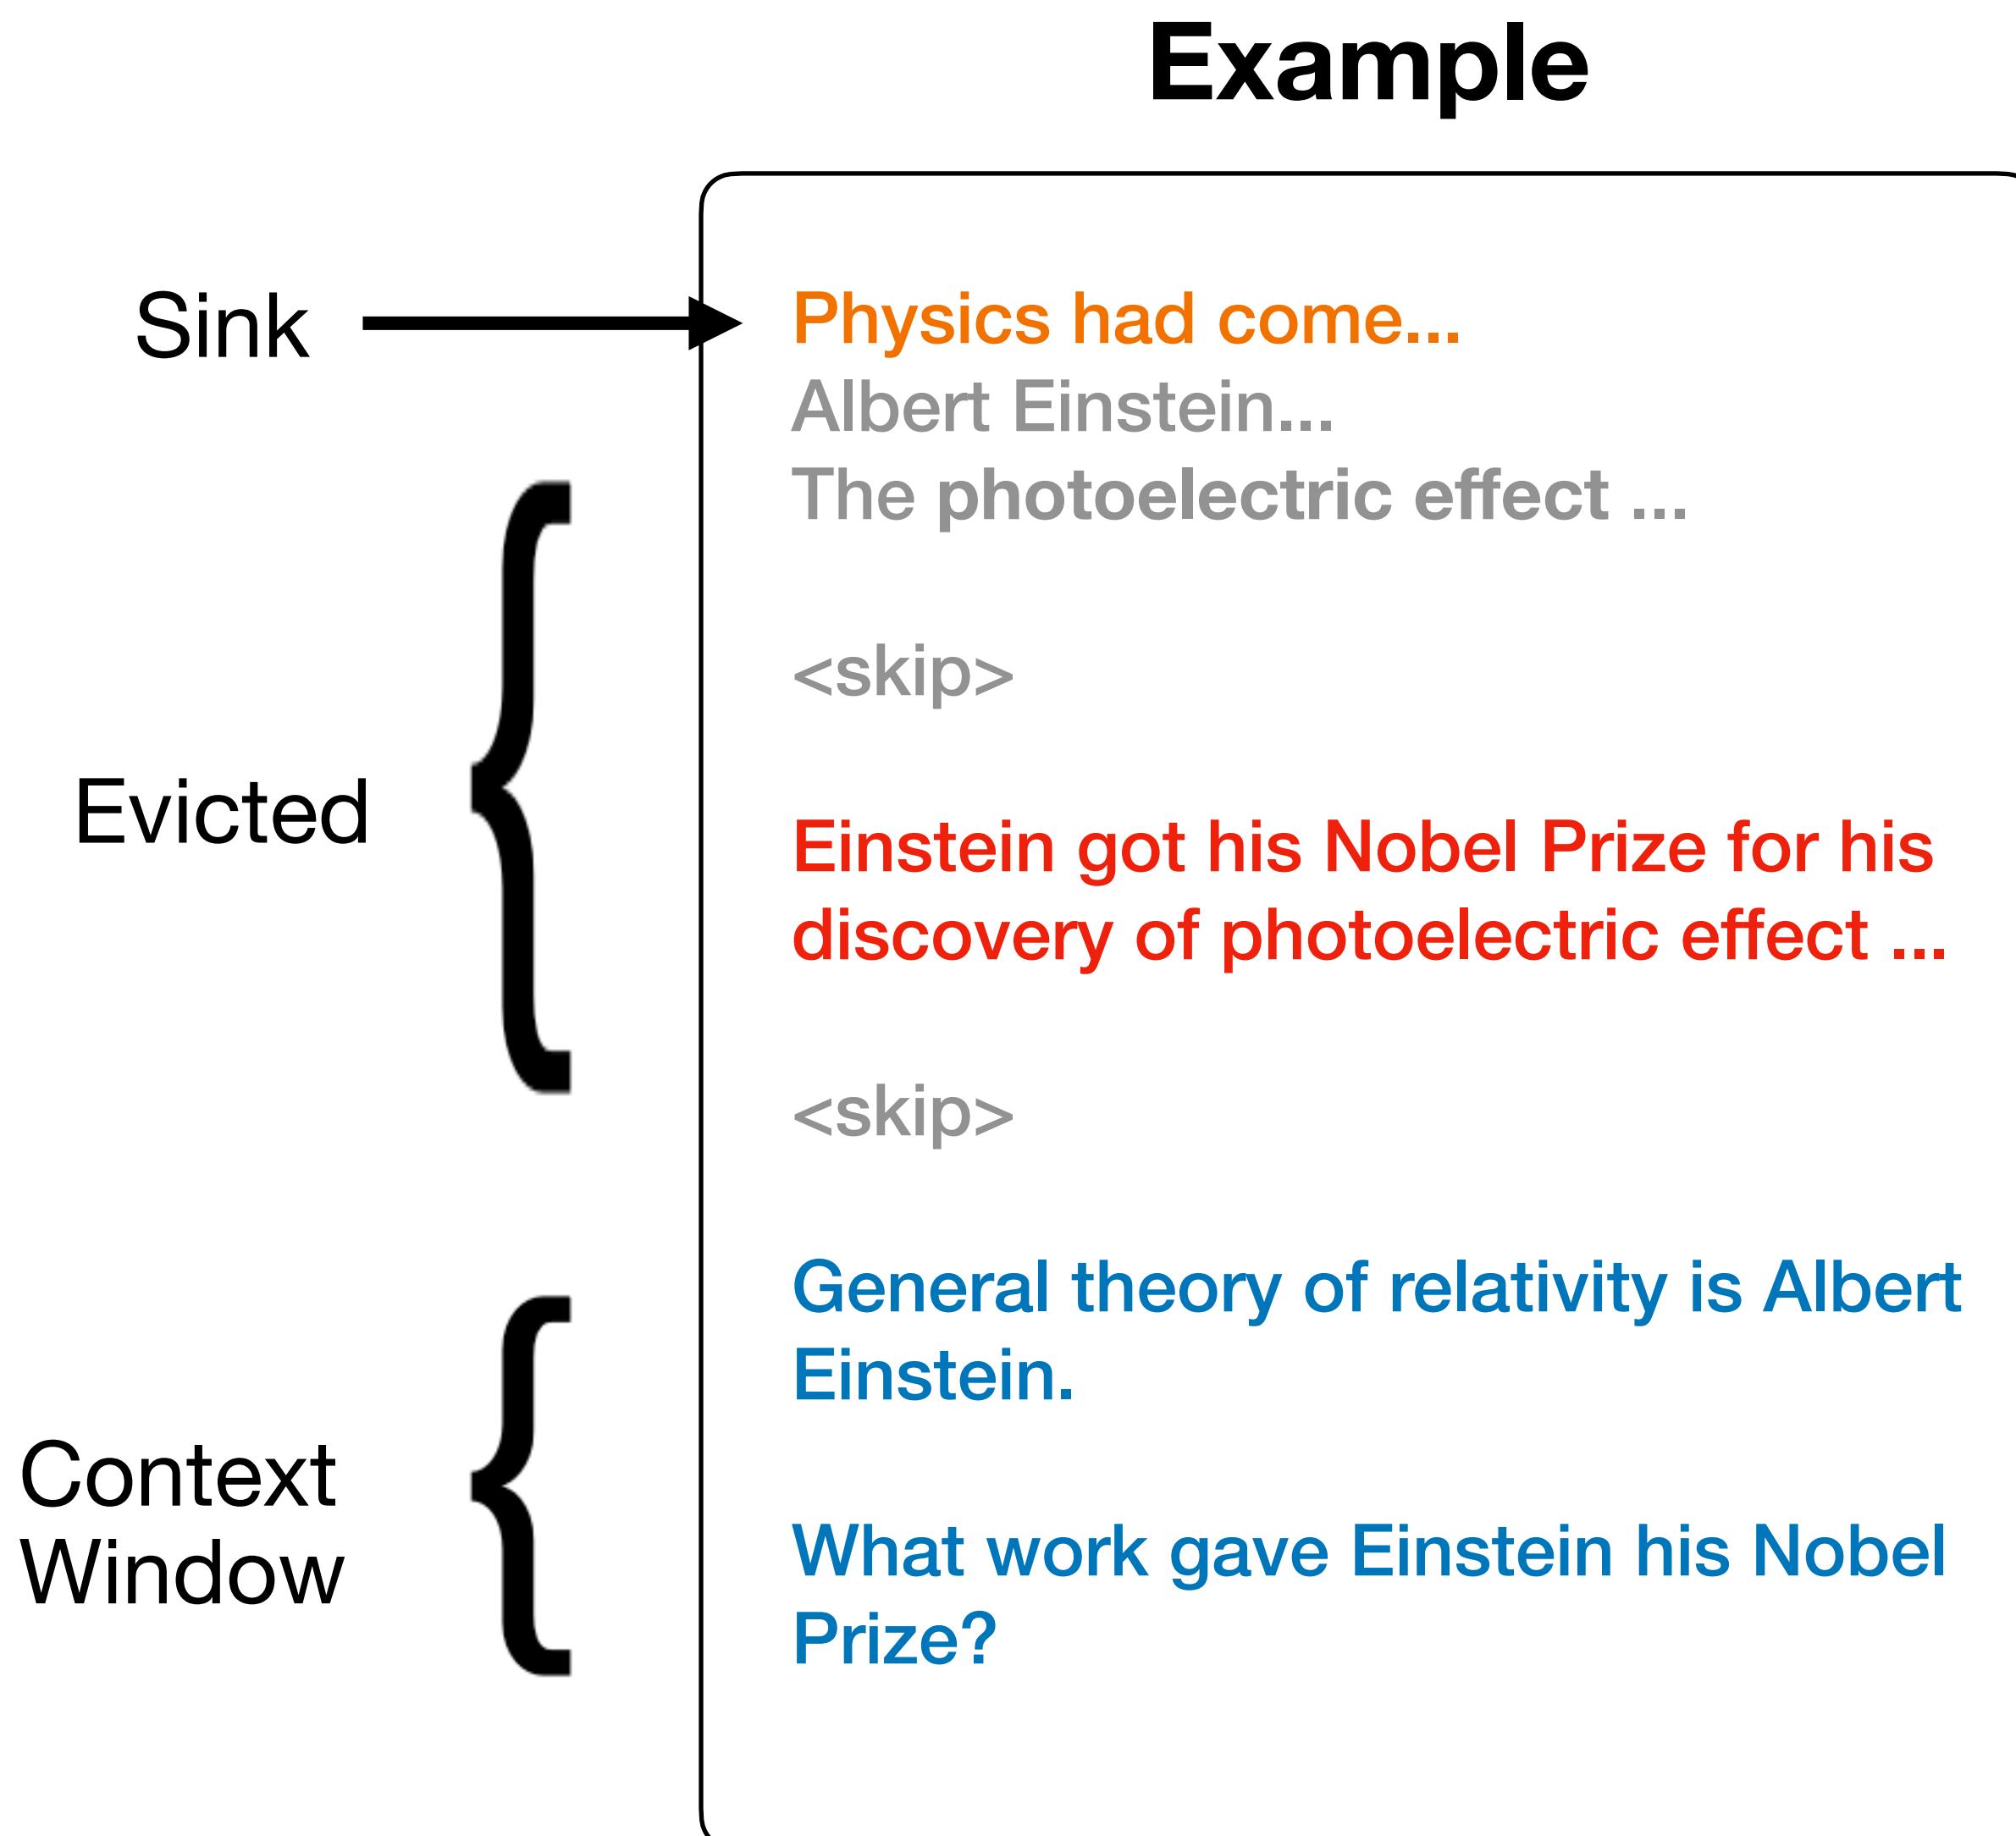
Example Prompt B

In this excerpt, answer this question.
Why was John absent yesterday?

Today was ...
For some reason **he fell sick** ...
So that was ...

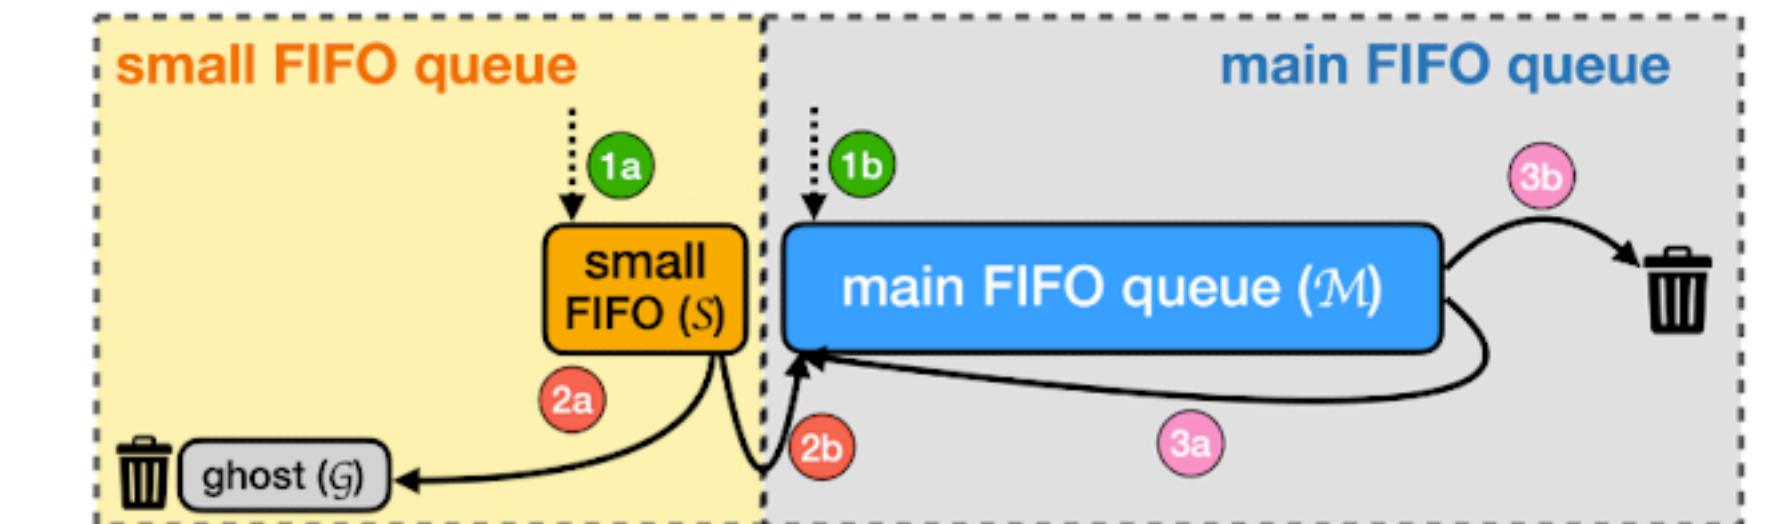
Problems with AttentionSinks

- Highly referenced tokens can exist at the middle.
- AttentionSinks and Context Windows **aggressively removes tokens at the middle** of the prompt and generated tokens.



Insights from Current Hardware-side Research

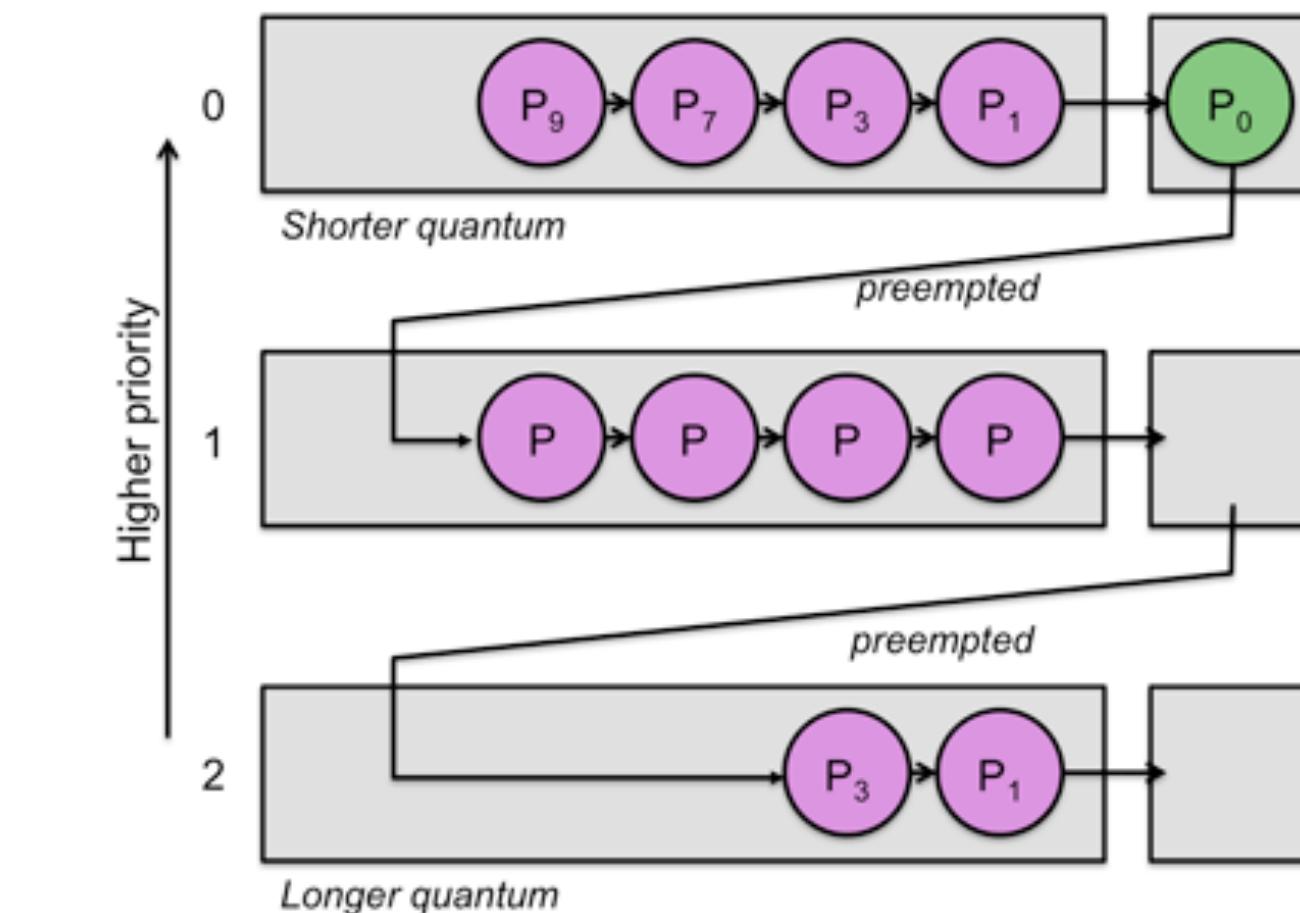
- Hierarchical-based policies are popular in hardware.
- **S3-FIFO (SOSP 2023):**
 - filter-out cold objects, keep hot objects
 - Hierarchy between cold and hot
- **MLFQ (Undergrad OS):**
 - hierarchy based on CPU usage of a program
 - Prioritize programs that use the CPU less



Insert: if not in ghost, insert to small **1a**, else insert to main **1b**

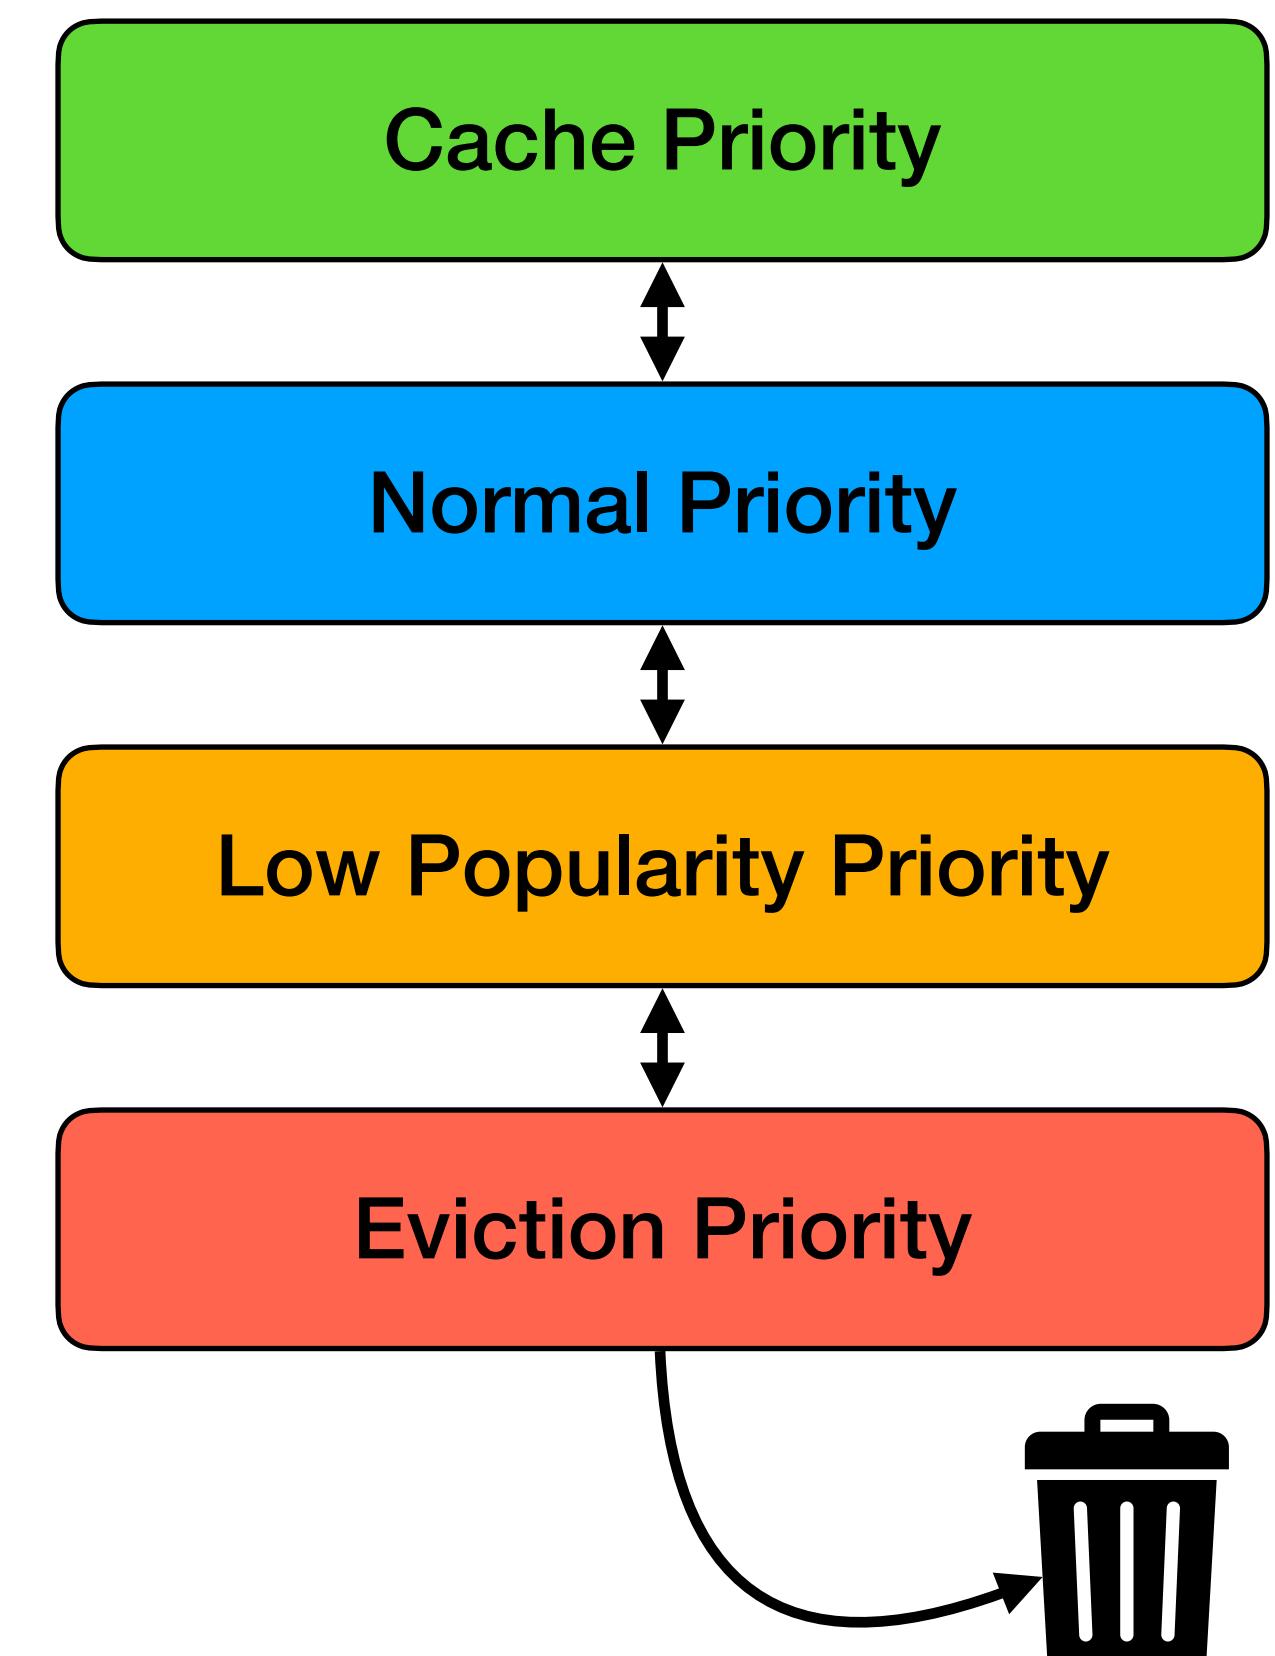
Evict (small): if not visited, insert to ghost **2a**, else insert to main **2b**

Evict (main): if visited, insert back **3a**, else evict **3b**



Motivation

- Current approaches (aggressively remove middle parts, or evict based on a small part of the prompt) are static.
- Hierarchical approaches allow for dynamic changes in the token access behavior.
 - Tokens that are **highly referenced** are placed at a **high caching priority**.
 - **Unpopular tokens** move toward the **eviction priority**.



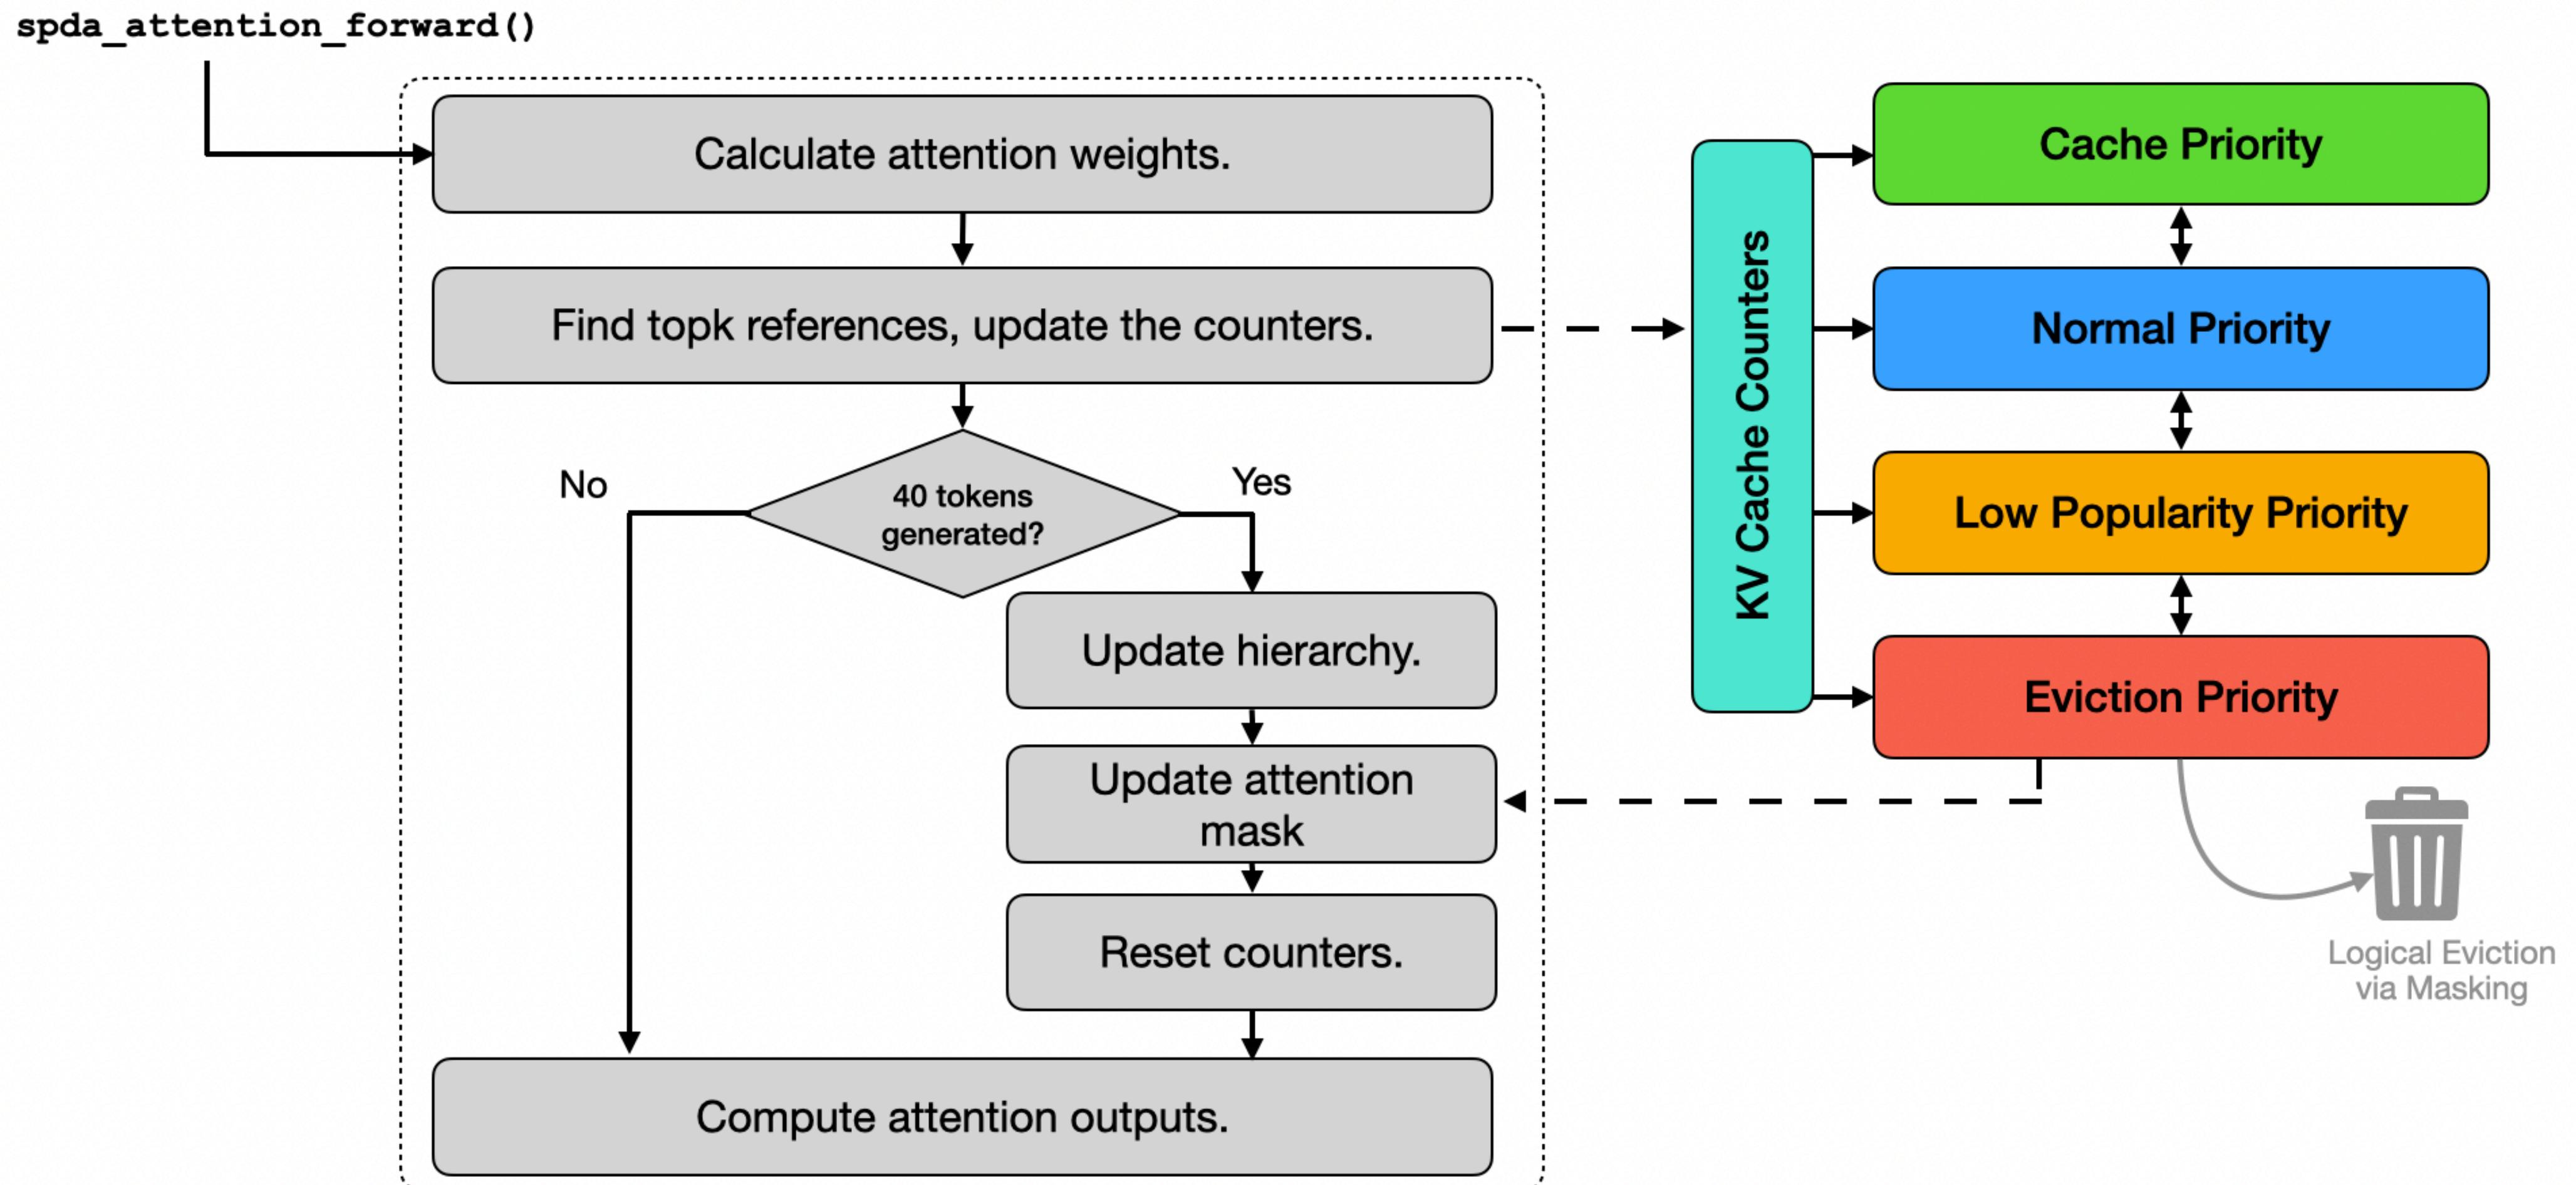
Implementation Setup

- **Interface:** Using HuggingFace
- **LLM:** Llama-2-7b-QNA-Tuned
- **Evaluation Dataset:** Open-source IELTS Reading Comprehension Test
- **Methodology:** Inject logical masking-based hierarchical KV cache management at `spda_attention_forward()` to emulate eviction.
 - Apparently, Pytorch does not have `explicit free()`.
 - llama.cpp was an option but for research purposes, logical eviction is sufficient.

Popularity Definition

- A token is considered popular if its attention weight often places it in the topk keys.
- Based on the number of times a token is part of topk within a profiling window, we place it on its corresponding hierarchy level:
 - **Cache Priority:** 10+ counts
 - **Normal Priority:** 5-9 counts
 - **Low Popularity Priority:** 1-4 counts
 - **Eviction Priority:** 0 counts
- The profiling is done every 40 tokens generated.
- Demotion by one level is performed if it fails to achieve required counts. Else, promotion.

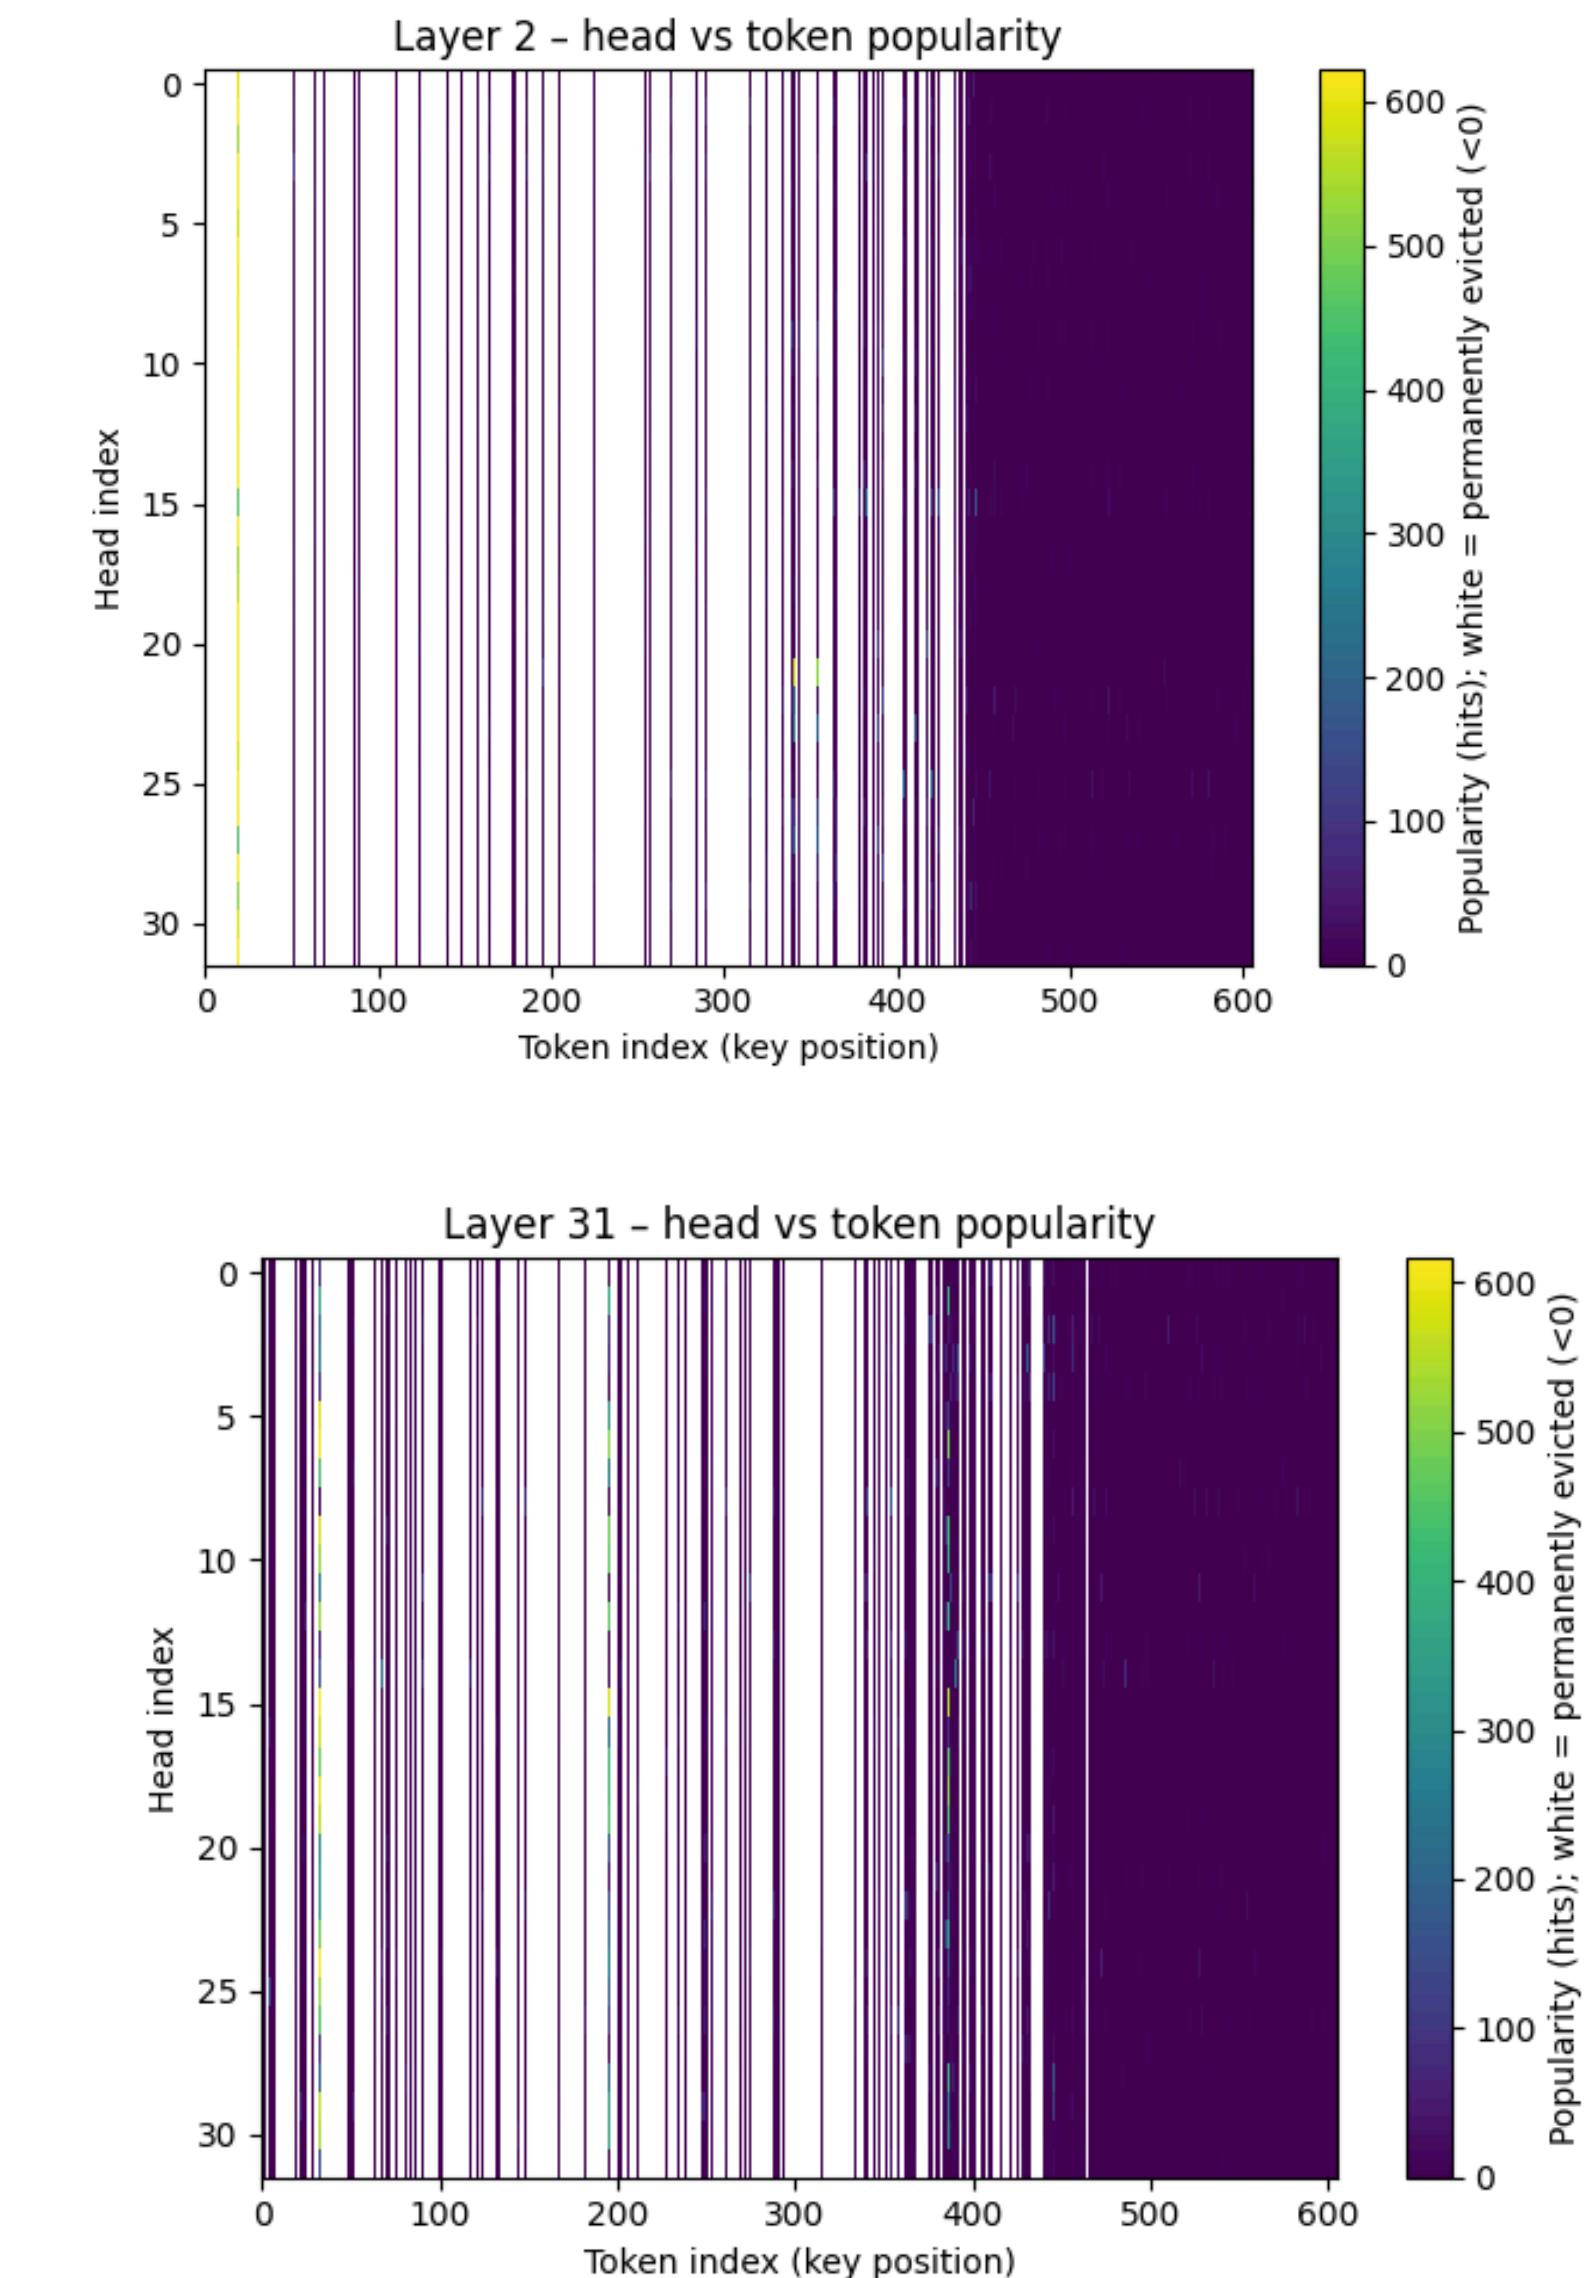
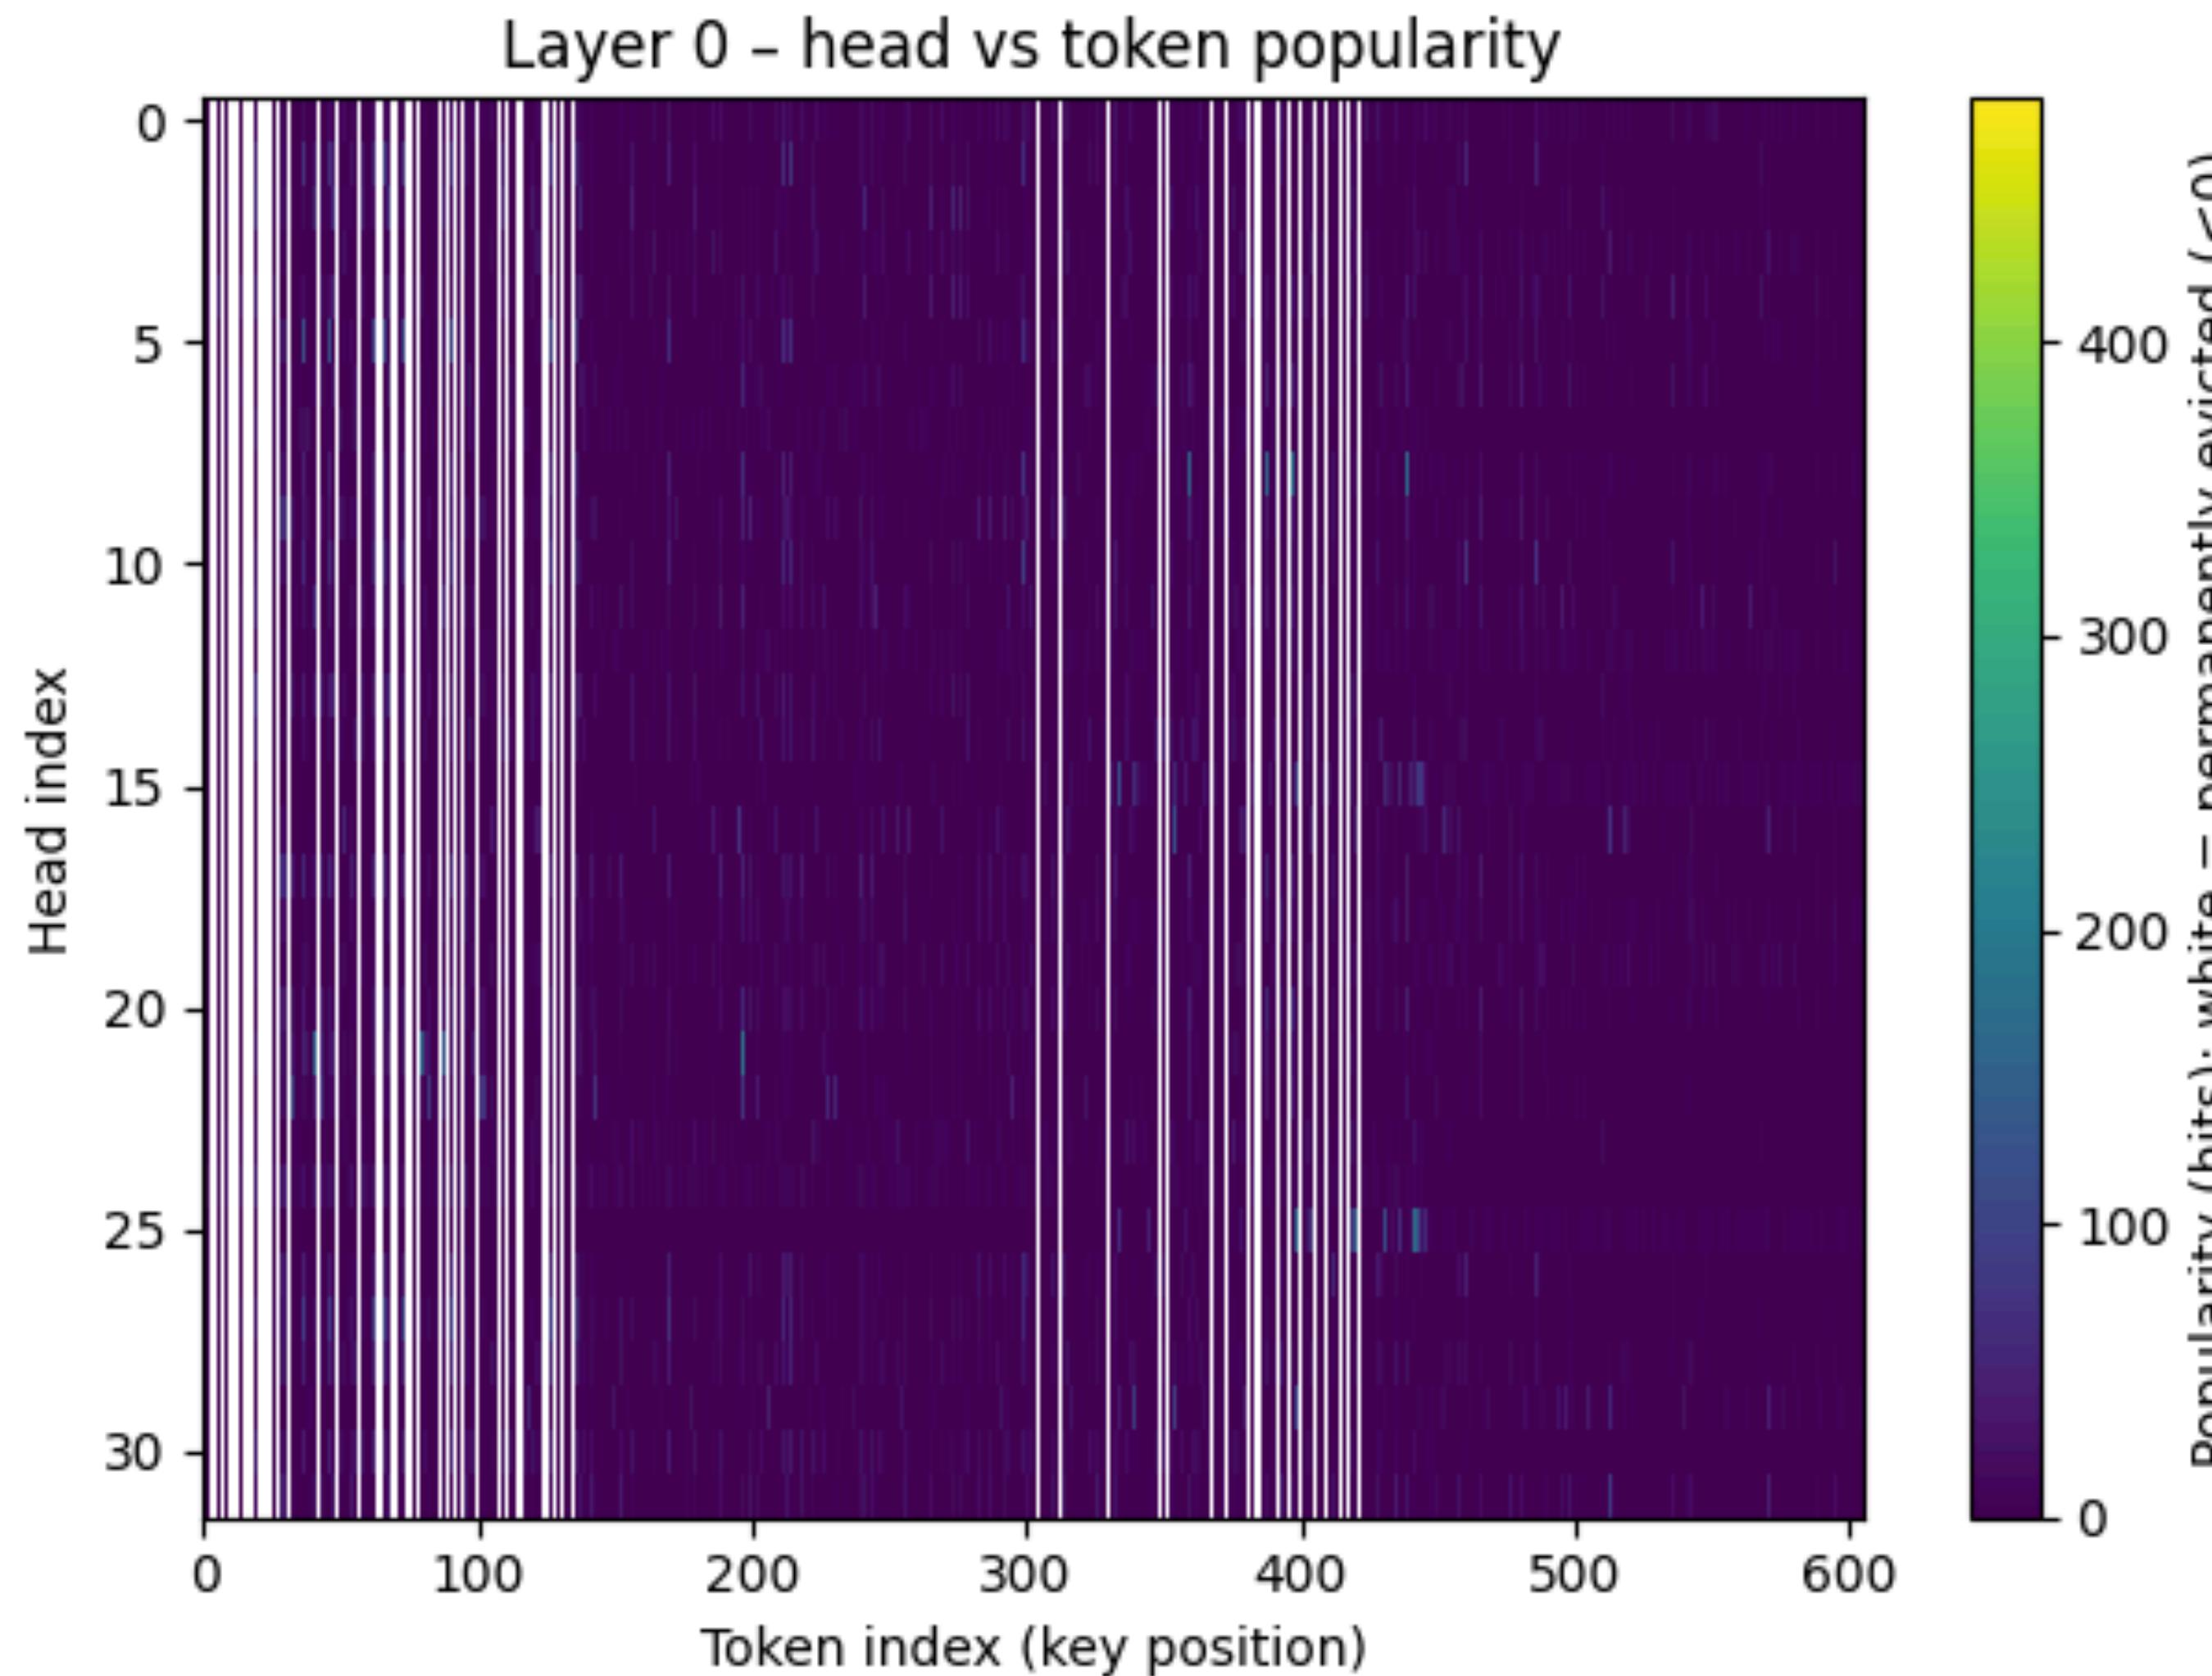
Policy Implementation



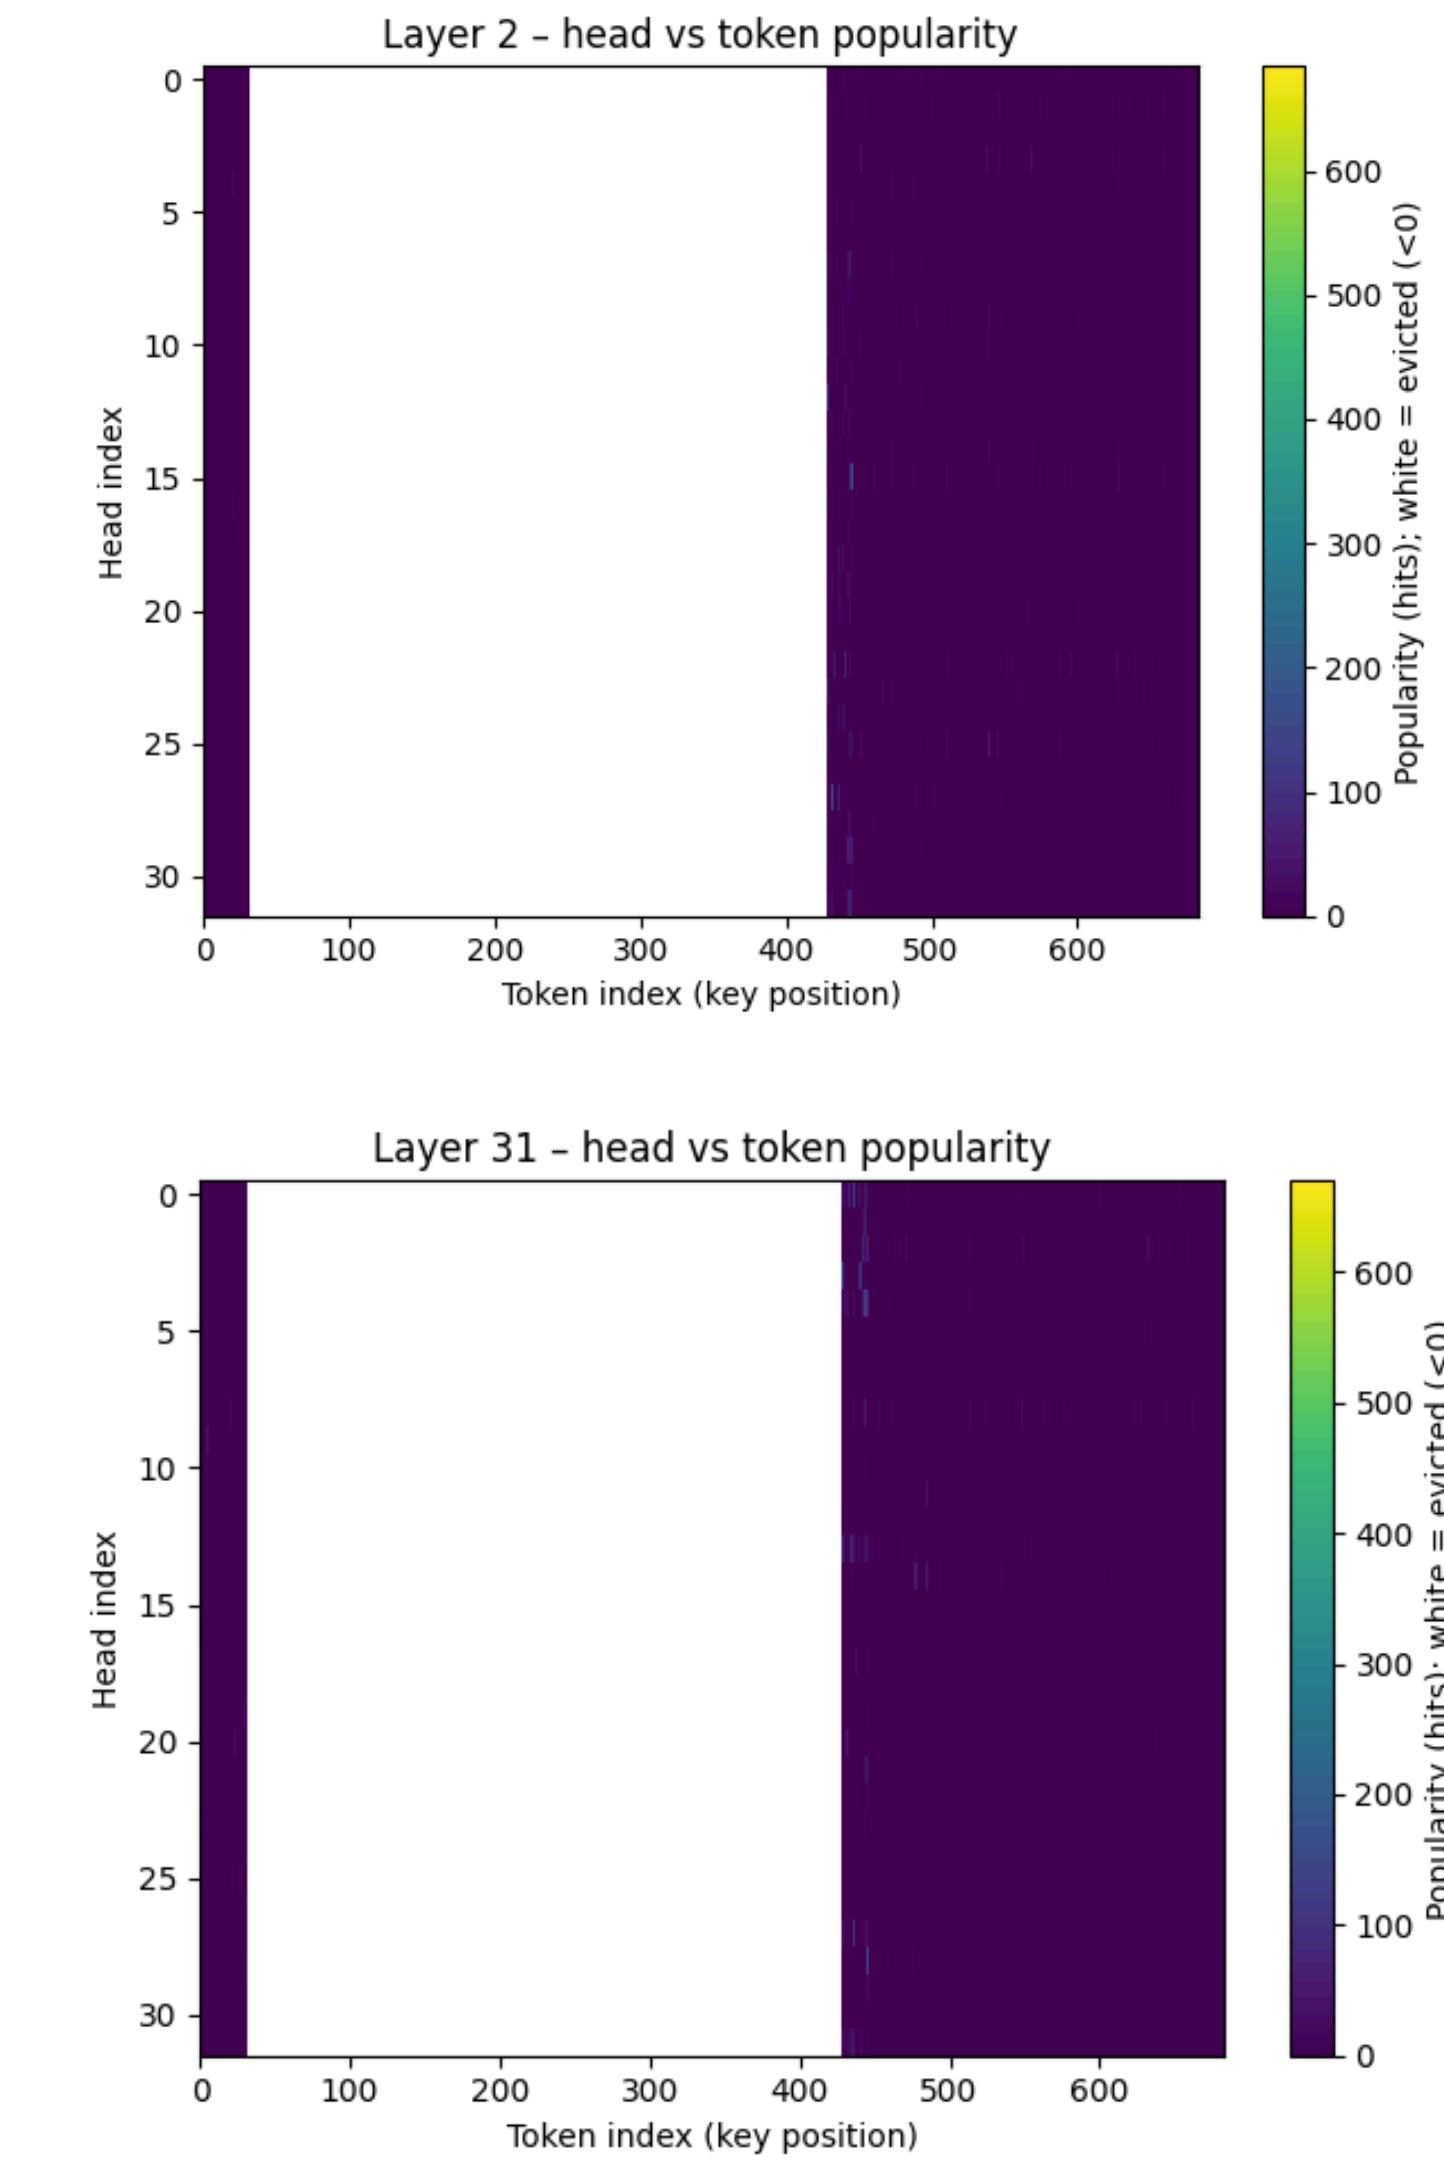
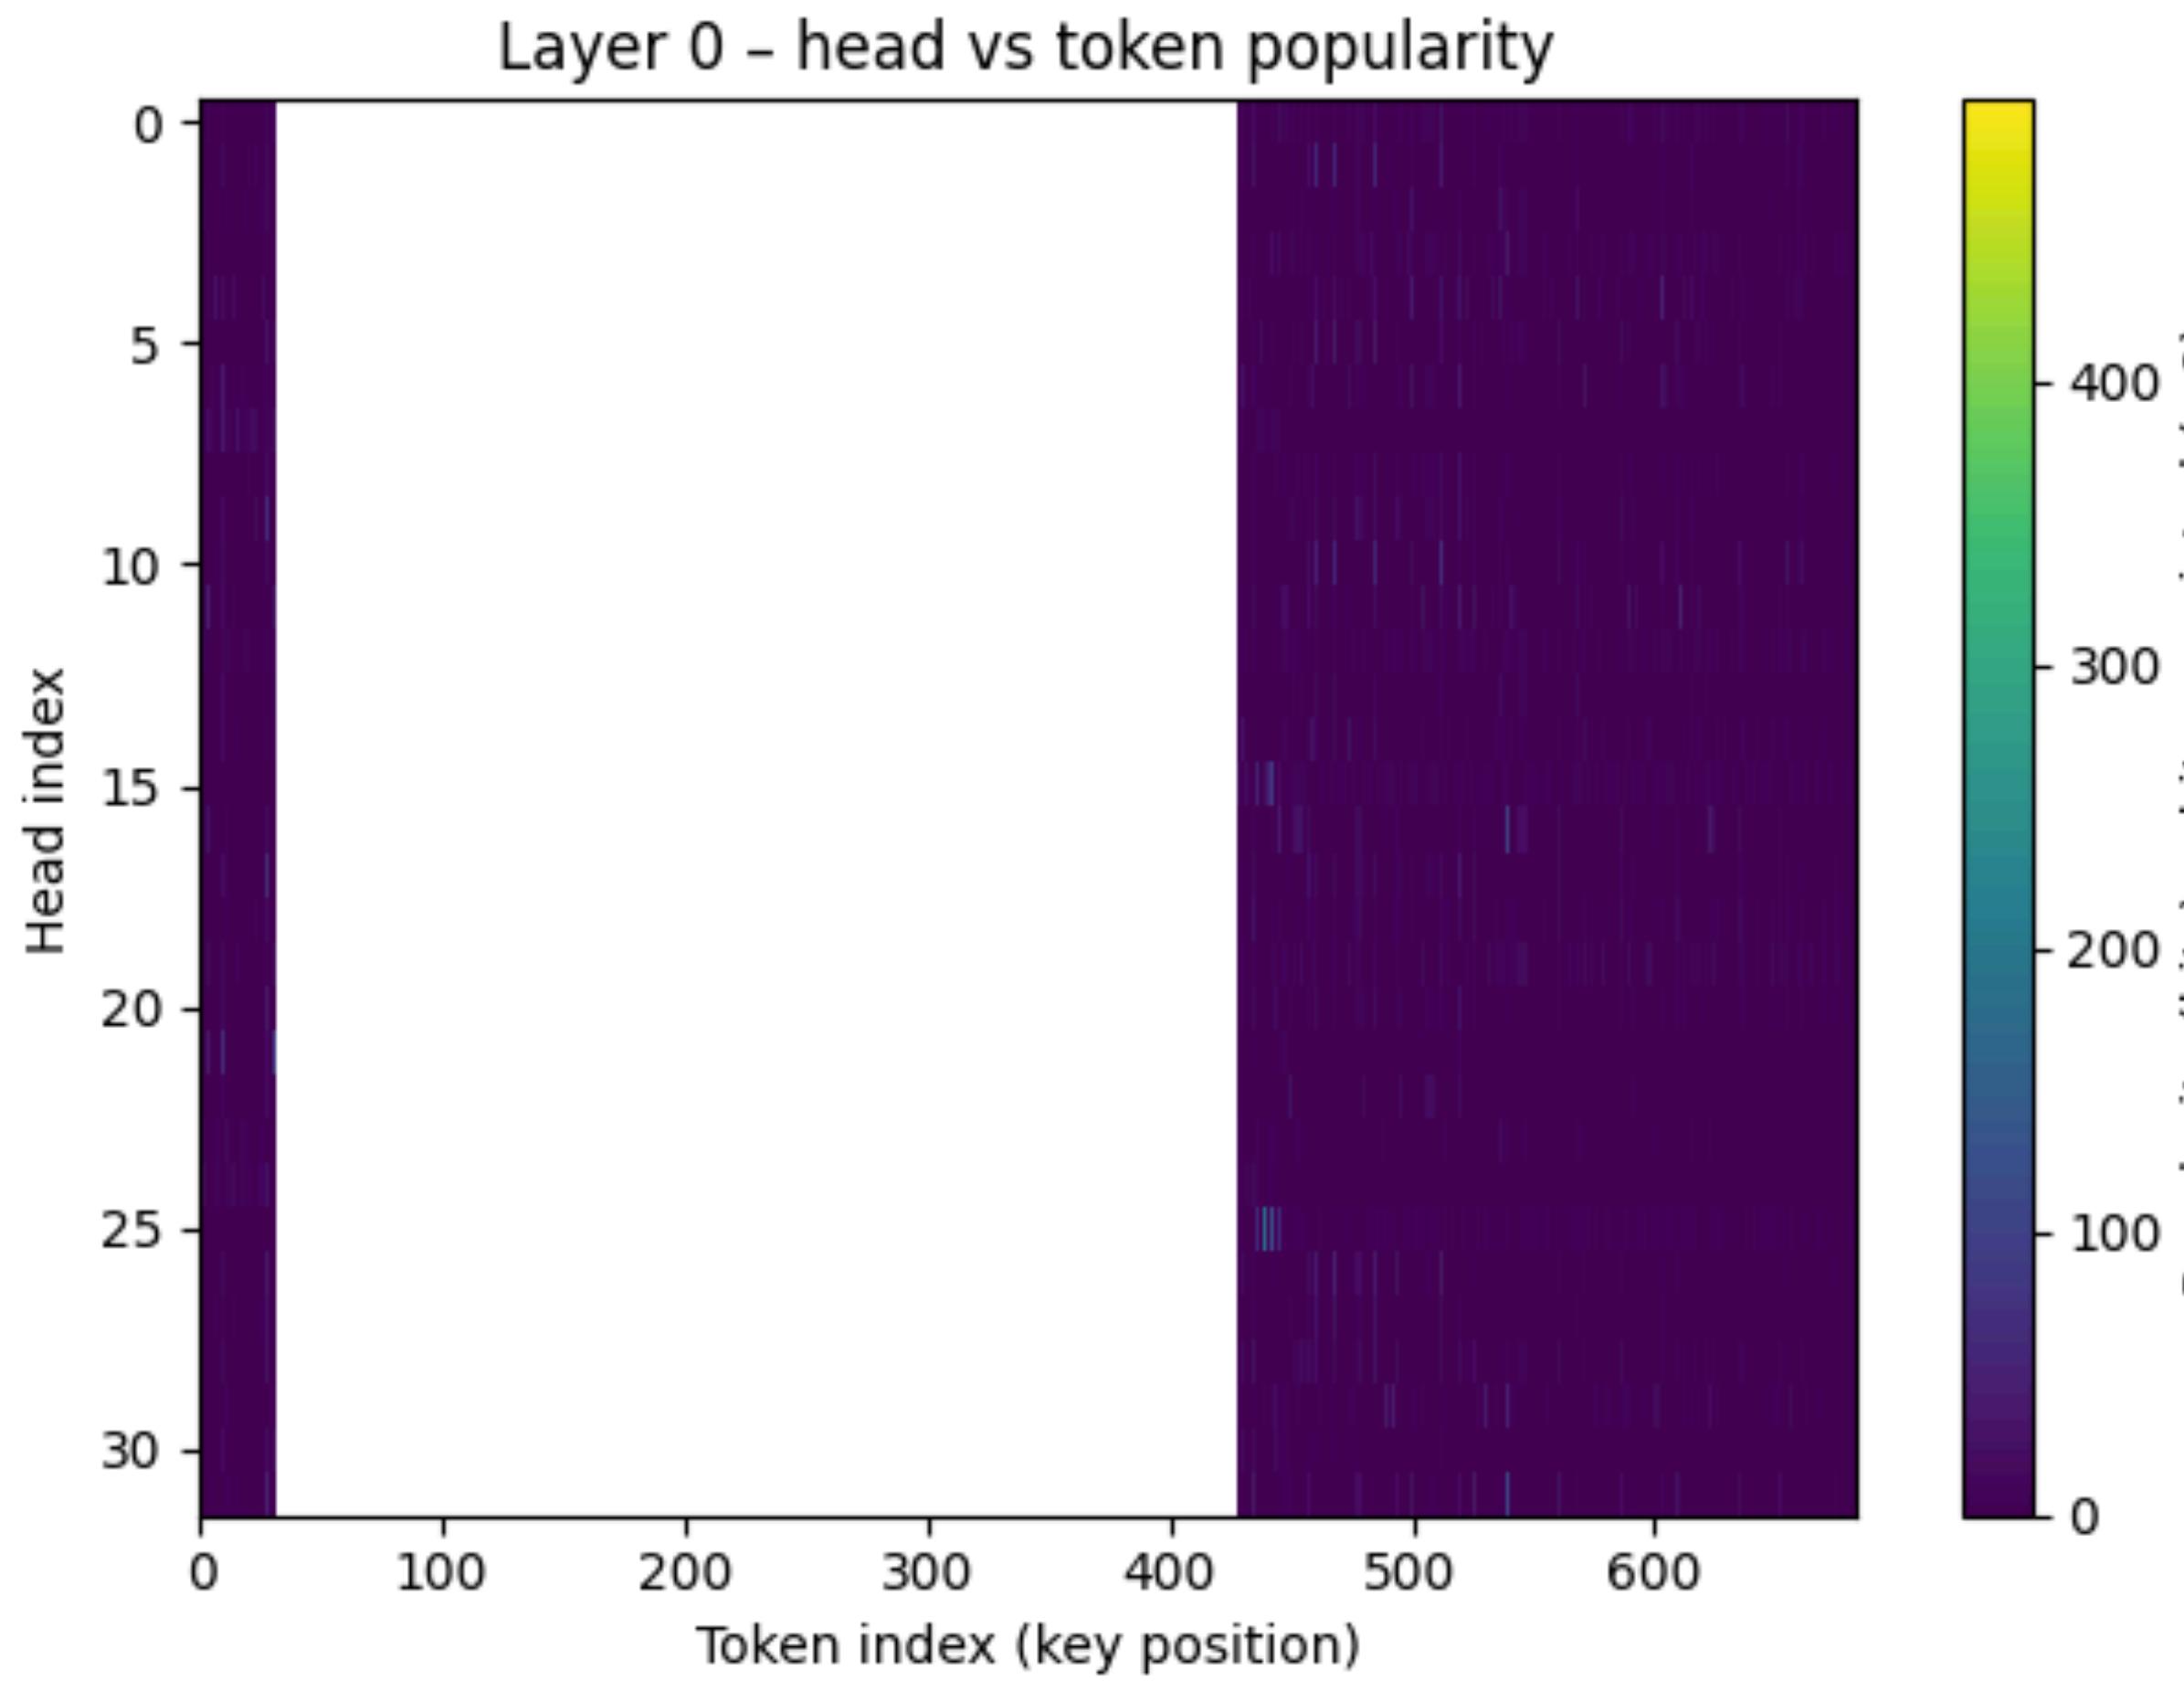
Preliminary Experiments

- **Prompts to test Needle-in-Haystack capability, four questions each:**
 - True or False, Fill in the Blanks, Multiple choice question, Matching Type
- **Position of the question:**
 - Before the passage
 - After the passage
- **Metric:** Quantitatively check if the LLM can answer correctly.
- **Comparison:**
 - No Eviction Baseline
 - Attention Sink (32 attention sink tokens, 256-token context window size)

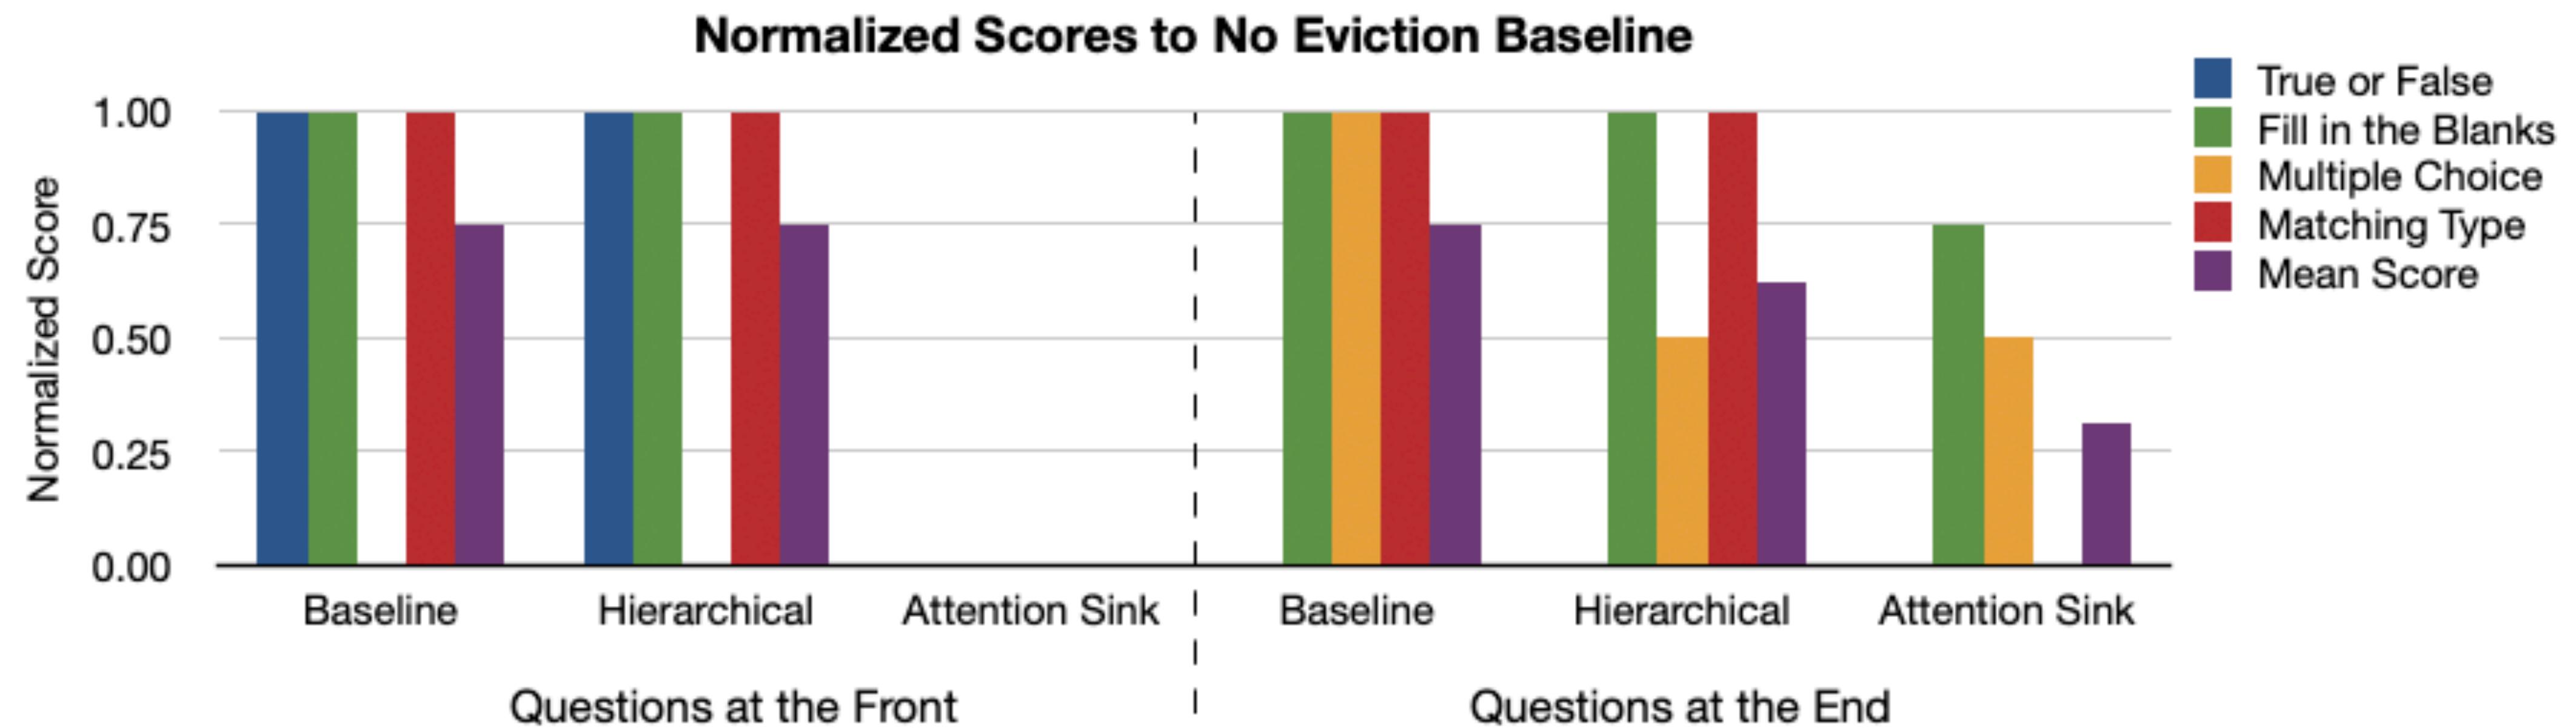
Popularity Plots for Hierarchical Policy



Popularity Plots for AttentionSink



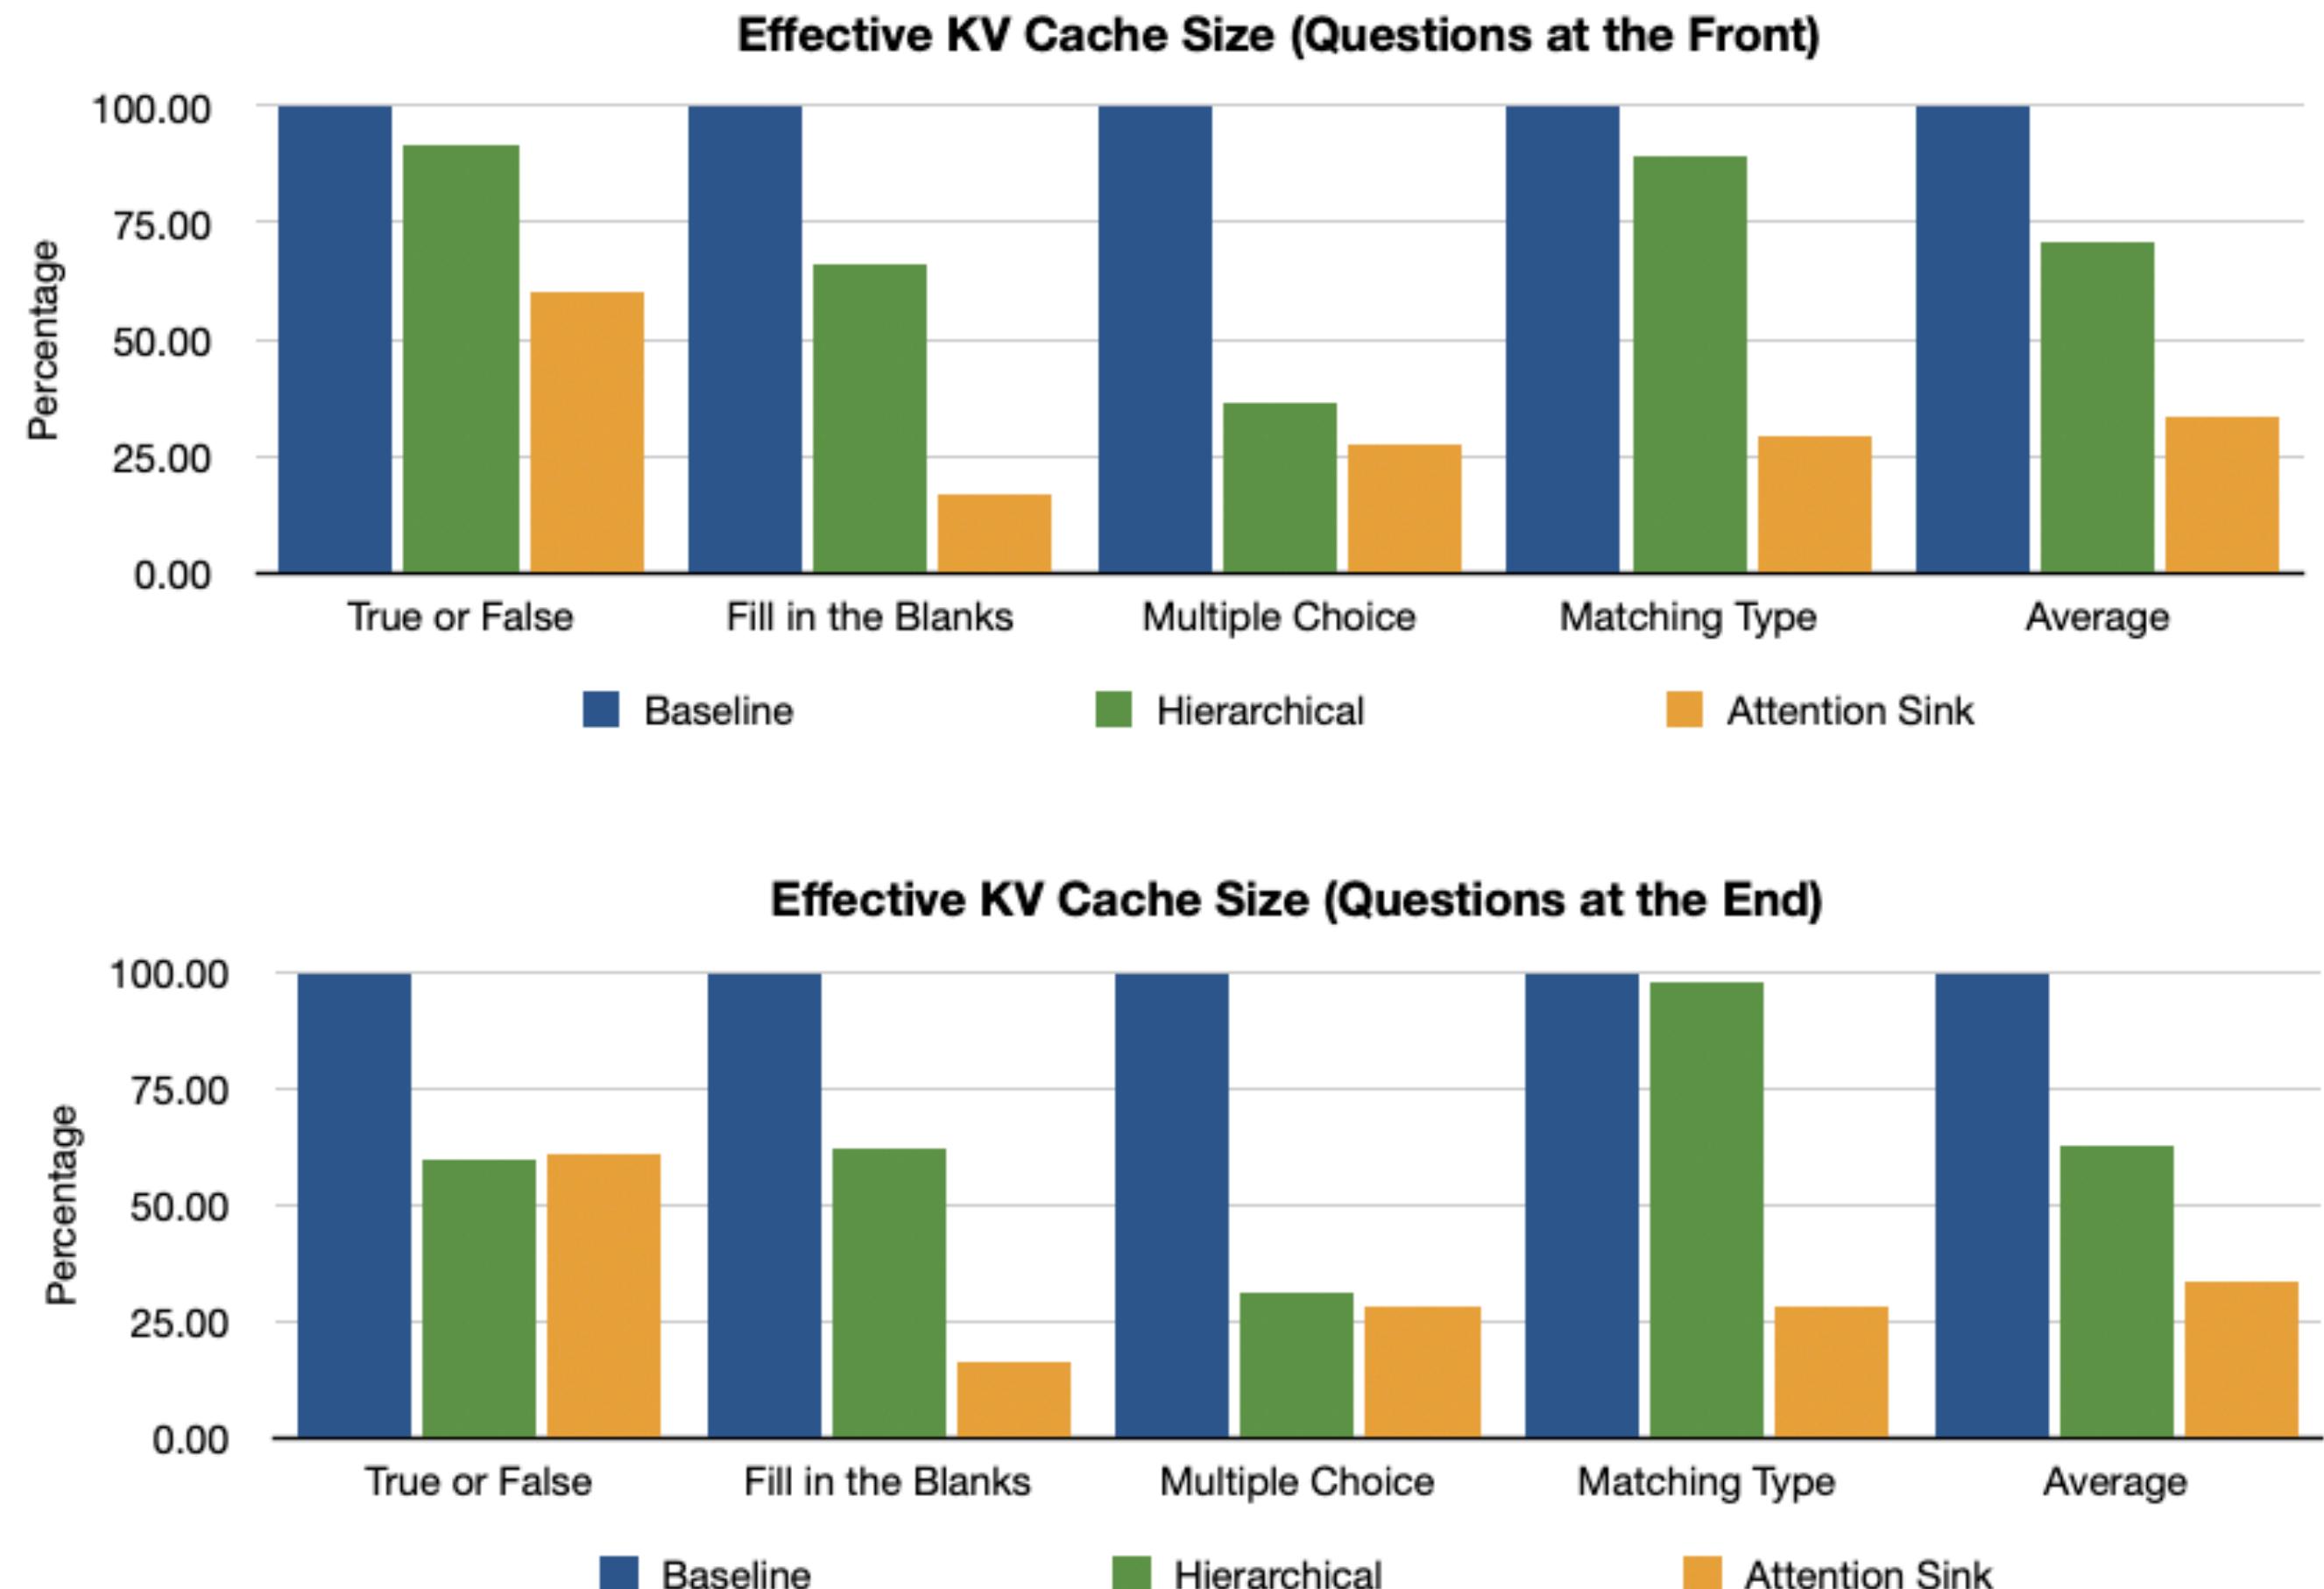
Results



- Hierarchical Policy shows relatively similar scores with Baseline.
- Attention Sink suffers significantly when questions are at the front.
 - Creates its own questions, or wrongly answers one question.

Results

- Hierarchical Policy shows diverse KV cache sizes which demonstrates some its ability to dynamically adapt to diverse workloads.
- Attention Sink aggressively drops all tokens in the middle.



Discussions

- Hierarchical KV Cache Compression Policy performs similarly with baseline with no eviction, demonstrating effectiveness.
- Effective KV Cache size varies among different types of questions which shows inherent ability to dynamically adapt to token popularity patterns.
- Demonstrates weakness of static and aggressive KV Cache compression techniques.
- This compression technique can be a stepping stone for a heterogenous memory system:
 - Highly popular KV entries placed in GPU memory
 - Less popular KV cache entries evicted to CPU memory